



## Evaluating Score-Investing Methodologies: A Systematic Review of Tweenvest's Algorithm for Long-Term Stock Investing Using Descriptive Analytics and Predictive Modeling.

Trabajo Fin de Grado Integrado  
Grado en Administración y Dirección de Empresas  
Grado en Ingeniería de Tecnologías y Servicios de Telecomunicación

**Author:** Carlos Eduardo Domínguez Martínez

**Tutor:** Joaquín Martínez Minaya (GADE)

**Tutor:** Alberto Albiol Colmener (GITST)

**Extern Co-Tutor:** José Tatay Sangüesa (Tweenvest)

**Course:** 2024-2025

Valencia, June 2025

## **Abstract**

This study investigates the effectiveness of a factor-investing methodology developed by Tweenvest, leveraging a proprietary algorithm grounded in fundamental financial analysis. The algorithm scores companies across four key factors: Quality, Growth, Value, and Dividends.

The research aims to evaluate the profitability of investment strategies based on these scores over multiple profitability horizons from 1 month to five years. A comprehensive dataset was constructed, integrating factor scores with additional variables such as sector and geographic region, standardized for currency and timeframes.

Statistical analyses will explore relationships between these factors and returns, identifying optimal investment periods. Subsequently, predictive modeling—including econometric regressions, time series models, and neural networks—will be applied to assess the translation of these factors into market performance.

The study incorporates an interdisciplinary approach, combining financial theory and econometrics with advanced programming, data engineering, and artificial intelligence. This integration bridges Business Management and Telecommunications Engineering, offering insights into the practical application of Tweenvest's scoring algorithm and contributing to the advancement of financial technology analytics.

To my parents.



# Contents

<b>I Main Report</b>	<b>1</b>
<b>1 Introduction and Objectives</b>	<b>3</b>
1.1 Fundamental Analysis Investing . . . . .	3
1.2 Problems With the Current Approach . . . . .	3
1.2.1 Unquantified Variables . . . . .	3
1.2.2 Emergent Effects in Complex Systems . . . . .	4
1.2.3 Linearity and Stationarity . . . . .	4
1.3 Objectives . . . . .	5
1.4 Variables for the Analysis . . . . .	5
1.4.1 Response Variables . . . . .	5
1.4.2 Explanatory Variables . . . . .	6
<b>2 Methodology &amp; Theoretical Framework</b>	<b>7</b>
2.1 Tweenvest's Scores . . . . .	7
2.1.1 Quality . . . . .	7
2.1.2 Growth . . . . .	10
2.1.3 Value . . . . .	11
2.1.4 Dividend . . . . .	12
2.2 Code practices . . . . .	13
2.3 Preprocessing . . . . .	14
2.3.1 Data Consistency . . . . .	14
2.3.2 Data Transformation . . . . .	15
2.4 Outlier Detection . . . . .	16
2.4.1 Inter Quartile Range . . . . .	16
2.4.2 Single Vector Machine . . . . .	17
2.4.3 Isolation Forest . . . . .	18
2.4.4 Local Outlier Factor . . . . .	19
2.4.5 Multi-Criteria Outlier Detection . . . . .	19
2.5 Predictive Models . . . . .	20
2.5.1 Regression Models . . . . .	20
2.5.2 Time Series . . . . .	21
2.5.3 Neural Networks . . . . .	22
2.6 Back-testing . . . . .	22
2.6.1 Performance . . . . .	23
2.6.2 Consistency and Overfitting . . . . .	23
2.6.3 Robustness of Results . . . . .	23
<b>3 Development</b>	<b>25</b>
3.1 Architecture Enhancement . . . . .	25
3.1.1 Updating Tweenvest Code . . . . .	25
3.1.2 Designing new Jobs . . . . .	27

3.1.3	Telemetry Tracing . . . . .	28
3.1.4	Dedicated Server Deployment . . . . .	30
3.2	Dataset Creation . . . . .	30
3.2.1	Aggregating the data . . . . .	30
3.2.2	Datasets Construction . . . . .	31
3.2.3	Preprocessing . . . . .	32
3.2.4	Outlier Detection . . . . .	32
3.2.5	Summary . . . . .	37
<b>4</b>	<b>Results</b>	<b>39</b>
4.1	Descriptive Analysis . . . . .	39
4.1.1	Variables Distributions . . . . .	39
4.1.2	Correlations . . . . .	42
4.1.3	Data clusters . . . . .	44
4.2	Predictive Models . . . . .	45
4.2.1	Linear Regression . . . . .	45
4.2.2	The problem with ARIMA . . . . .	45
4.2.3	GAM . . . . .	45
4.2.4	Neural Network . . . . .	45
4.3	Back-testing . . . . .	45
<b>5</b>	<b>Discussion</b>	<b>47</b>
<b>6</b>	<b>Conclusion</b>	<b>49</b>
<b>Bibliography</b>		<b>51</b>
<b>II</b>	<b>Appendix</b>	<b>53</b>
<b>7</b>	<b>Code Listings</b>	<b>55</b>
7.1	Historical Scores Job Implementation . . . . .	55
7.1.1	Enqueueing Historical Scores Job . . . . .	55
7.1.2	Calculate Index Scores Data . . . . .	56
7.1.3	Calculate Stocks Scores Task . . . . .	58
7.2	Workers Management . . . . .	59
7.2.1	Tweenvest's Production Server . . . . .	59
7.2.2	Personal Development Server . . . . .	61
7.3	Data Export Job Implementation . . . . .	61
7.4	Custom Python Functions for Data Analysis . . . . .	65
7.5	Data Preprocessing . . . . .	66
7.5.1	Data Cleaning . . . . .	66
7.5.2	Outlier Detection . . . . .	67
7.6	Extra Images . . . . .	67
7.6.1	Data . . . . .	68
7.6.2	Distributions . . . . .	68
7.6.3	Correlations . . . . .	70

# List of Figures

2.1 Tweenvest's Quality Score . . . . .	10
2.2 Tweenvest's Growth Score . . . . .	11
2.3 Tweenvest's Value Score . . . . .	12
2.4 Tweenvest's Dividend Score . . . . .	13
2.5 Yeo-Johnson Transformations examples . . . . .	16
2.6 Visualization of how Isolation Forest works. . . . .	18
2.7 Multi-Criteria Outlier Detection Method . . . . .	20
3.1 General Tweenvest's Architecture . . . . .	25
3.2 Scores Calculations Code Schema . . . . .	26
3.3 <i>calculation_date</i> Field Propagation Schema . . . . .	26
3.4 Historical Jobs First Iteration Schema . . . . .	28
3.5 Historical Jobs Final Version Schema . . . . .	28
3.6 First Data Aggregation Schema . . . . .	30
3.7 Final Data Aggregation Schema . . . . .	31
3.8 Past Dataset . . . . .	31
3.9 Future Dataset . . . . .	31
3.10 Small Data Future - IF (5% tolerance) . . . . .	33
3.11 Small Data Future - SVM (5% tolerance) . . . . .	34
3.12 Small Data Future - LOF (5% tolerance) . . . . .	35
3.13 Small Data Future - Multi-Criteria (20% tolerance) . . . . .	36
4.1 Big Data Future - Scores . . . . .	39
4.2 Big Data Future - Scores Deep . . . . .	39
4.3 EMMI - Euribor Interest Rates . . . . .	40
4.4 Small Data Future - Dummies . . . . .	41
4.5 Big Data Future IQR - Profits Boxplot . . . . .	41
4.6 Big Data Future IQR - Profits . . . . .	41
4.7 Big Data Past - Correlations . . . . .	42
4.8 Big Data Future - Correlations . . . . .	42
4.9 Big Data Past IF - Correlations . . . . .	43
4.10 Big Data Future IF - Correlations . . . . .	43
4.11 Pairplot of Small Data Future - Multi . . . . .	44
7.1 Number of columns per dataset . . . . .	68
7.2 Big Data Past - Scores . . . . .	68
7.3 Big Data Past - Scores Deep . . . . .	68
7.4 Small Data Future - Scores . . . . .	69
7.5 Small Data Future - Scores Deep . . . . .	69
7.6 Big Data Future - Dummies . . . . .	69
7.7 Big Data Past - Dummies . . . . .	69

7.8	Big Data Future - Profits Boxplot . . . . .	69
7.9	Big Data Future - Profits . . . . .	69
7.10	Small Data Future USA . . . . .	70
7.11	Small Data Future USA - Multi . . . . .	70
7.12	Small Data Future EU . . . . .	70
7.13	Small Data Future EU Multi . . . . .	70
7.14	Small Data Future LAT . . . . .	70
7.15	Small Data Future LAT - Multi . . . . .	70
7.16	Small Data Future AF . . . . .	71
7.17	Small Data Future AF - Multi . . . . .	71
7.18	Small Data Future AS . . . . .	71
7.19	Small Data Future AS - Multi . . . . .	71

# List of Tables

2.1	Pros and Cons of the Interquartile Rule . . . . .	17
2.2	Pros and Cons of One-Class SVM for Outlier Detection . . . . .	17
2.3	Pros and Cons of the Isolation Forest Method for Outlier Detection . . . . .	19
2.4	Pros and Cons of the Local Outlier Factor (LOF) Method . . . . .	19
3.1	Tweenvest Server Specifications . . . . .	30
3.2	Hetzner Server Specifications . . . . .	30
3.3	Missing values (%) per column for Small Data Future. . . . .	32
3.4	Rows per Dataset . . . . .	37

**Part I**

**Main Report**



# Chapter 1

## Introduction and Objectives

### 1.1 Fundamental Analysis Investing

Fundamental Analysis is a methodology used to evaluate the intrinsic value of a company, asset, or market by analyzing various economic, financial, qualitative and quantitative factors. Unlike technical analysis, which focuses on price movements and chart patterns, fundamental analysis seeks to determine an asset's "true value" to identify investment opportunities that may be undervalued or overvalued in the market.

This intrinsic value is defined by numerous experts and renowned investors—such as Benjamin Graham, Warren Buffett, and Pat Dorsey [7]—as a company's ability to adapt to an ever-changing environment, create value in the market, and establish barriers to entry for competitors, commonly referred to as *economic moats*.

Measuring these moats is challenging due to the difficulty of quantifying certain variables, such as brand strength and market influence. However, over the long term most literature show that these **intangible factors translate into tangible financial data**, reflected in a company's reports: balance sheet, income statement, and cash flow statement. When exposed to the public market share transform into the need to buy or sell the stocks, altering the companies stocks' profitability. So from now on, this economic moats will be called *alpha*, Sharpe [14].

This complexity is what made Tweenvest develop a series of scores to help users identify possible *alpha* in a company, based on the fundamental analysis principles.

### 1.2 Problems With the Current Approach

In the context of this thesis, we have to explain the need to make a systematic review of the scores to check whether the simplifications and hypotheses assumed by the financial consensus, and the way that they are implemented in Tweenvest's scores, truly show a company's ability to create *alpha* over time or not.

#### 1.2.1 Unquantified Variables

Since the current model only includes variables found in a company's financial statements, there is a significant amount of relevant information being left out, for example:

- The **perceived differentiation** of a product is one such element that may not be properly captured by traditional financial analysis. A trusted brand or strong reputation can

allow a company to charge premium prices and build customer loyalty, generating extraordinary long-term profits. For instance, brands like Tiffany or Rolex have high perceived value that justifies elevated prices—something that cannot be easily quantified using standard financial metrics such as profit margins or return on capital.

- Additionally, strategies such as **cost leadership** and offering products at lower prices can provide a significant competitive advantage that is not always directly reflected in financial data.
- Companies may also establish **barriers to entry** and **high switching costs** for customers, making it more difficult for them to move to competitors. These strategies may involve investments in technology, patents, or simply building long-term customer and employees relationships.

To properly measure these variables, a more exhaustive analysis of each company would be required, pulling from multiple secondary information sources such as news articles, conference transcripts, customer blogs, competitive product reviews, and more. This task is far more difficult to automate via code, as it would require multimodal AI techniques—thus, it will fall outside the scope of this work.

### 1.2.2 Emergent Effects in Complex Systems

In complex systems like financial markets, network effects emerge, adding noise to actual reliable data behind it. For example:

- **Herd behavior** is a network effect where investors tend to follow the actions of others rather than rely on their own analysis. This can amplify market movements, both upward and downward. Herd behavior can lead to speculative bubbles and abrupt corrections when collective expectations shift.
- **Feedback loops** are another effect where market participants' actions reinforce existing market behavior. For example, rising asset prices may attract more investors, which in turn drives prices even higher. This type of positive feedback can cause price escalation that becomes detached from underlying fundamentals. Conversely, negative feedback can occur during a sell-off, where falling prices trigger more selling, further accelerating the decline.
- **Macro-level influences** are another crucial factor where broader structural changes—such as political decisions, economic policies, technological shifts, or industry-wide trends—can significantly impact market behavior. These larger forces often operate beyond the scope of individual company analysis and can create ripple effects throughout the entire market ecosystem.

### 1.2.3 Linearity and Stationarity

#### Non-linearity

The relationships between financial variables are often non-linear, meaning that changes in one variable may not result in proportional changes in another.

- **Economies of scale** can create non-linear relationships between production volume and costs. As a company grows, it may experience decreasing marginal costs due to better resource utilization, bulk purchasing discounts, or spreading fixed costs over larger output.

- **Market saturation** can lead to diminishing returns on marketing spend or R&D investments. Initial investments might yield significant returns, but as the market becomes saturated, additional spending may produce smaller incremental benefits.
- **Competitive dynamics** can create threshold effects where small changes in market share or pricing can trigger significant shifts in competitive position or profitability.

## Trends

Traditionally, standard growth ratios have been used, based on the assumption that in competitive markets, when a new business model or opportunity emerges with above-average margins, entrepreneurs quickly move in to capitalize on it—eventually saturating the opportunity and driving margins back down to average levels over time. But what happens when there is a significant shift in trend?

Markets are "fluctuating entities", so static metrics can become problematic when underlying trends change. A clear example is the rise of artificial intelligence and the surge in stock prices of companies involved in the production and development of the necessary technologies.

## 1.3 Objectives

Given the limitations identified in the current approach, Tweenvest needs to **evaluate if the scores represent *alpha* for different time frames**, and **guide users how to interpret them based on different investment strategies**. Additionally, to propose possible changes to the current algorithm in order to present the results with more accuracy. To achieve this, we follow a structured, multi-step process:

1. **System Architecture Enhancement:** Modify Tweenvest's database architecture to enable historical score storage and retrieval and create the datasets for the analysis.
2. **Data Curation:** Clean and preprocess the collected data to ensure consistency and reliability, handling missing values and outliers appropriately.
3. **Exploratory Analysis:** Process and analyze the data to check the distributions and correlations between variables, looking for early on patterns to later compare and use in the modeling.
4. **Predictive Modeling:** Develop multiple predictive models to better understand the scores and their relations with other variables and time frames.
5. **Validation & Benchmarking:** Back-test the algorithms to check their performance and consistency to see if the scores generate positive *alpha* over time.

## 1.4 Variables for the Analysis

For doing so, we are going to look for correlations between the profits of a company with different scores and other categorical variables.

### 1.4.1 Response Variables

To help users understand how to use the scores depending on their investment horizon, we used different profits that from now one will be represented as  $P_i$  for each time horizon:  $i \in \{1 \text{ month}; 3 \text{ months}; 6 \text{ months}; 1 \text{ year}; 2 \text{ years}; 5 \text{ years}\}$

The profits are defined as:

$$R_i^n = \frac{(P_{i+n} + D_{acc}) - P_i}{P_i} \quad (1.1)$$

Here,  $R_i^n$  represents the total return achieved over a period of  $n$  months starting from time  $i$ . This return accounts for both the price appreciation and the dividends received during the specified time frame. Specifically,  $P_i$  denotes the price at time  $i$ , while  $D_{acc}$  is the sum of all dividends accumulated from time  $i$  to  $i + n$ .

## 1.4.2 Explanatory Variables

To relate the profits with the scores, we need to use different variables:

- **Main variables**, these are numerical variables from **Tweenvest's Scores**, that are in the range of 0 to 100:  $S_j$  where  $j \in \{ \text{Quality}; \text{Value}; \text{Dividend}; \text{Growth} \}$
- **Dummy variables**, these are boolean variables that allow us to tag the companies by:
  - **Industry**:  $I_k$  where  $k \in \{ \text{Financials}; \text{Healthcare}; \text{etc.} \}$
  - **Region**:  $R_l$  where  $l \in \{ \text{Europe}; \text{North America}; \text{etc.} \}$
- **Market Variables**, these are numeric variables that represent the companies "size" using the: **Market Cap** and **Volume**.

# Chapter 2

# Methodology & Theoretical Framework

## 2.1 Tweenvest's Scores

The **long-term advantage** explained earlier in the finance is used in many indexes such as MSCI [12], IBEX-35, etc to achieve higher returns than average. So following this approach, Tweenvest developed a series of scores to **help users identify possible alpha in a company**. Upon these scores, we can highlight four –Quality, Growth, Value, and Dividend– that are strictly related to key focus areas that fundamental-based investors consider before making any decision, these areas are:

- Profitability
- Financial Health
- Predictability
- Consistent Growth
- Entrance Moment
- Dividends Payed

### 2.1.1 Quality

This Tweenvest's score is approached in a similar way that many successful investors would when analyzing the annual reports, by distinguishing on three main categories: **profitability**, **financial health**, and **predictability**. Each of them includes inside of them multiple financial ratios that take account of different relevant data within the company's reports.

To understand the complexity of this score, we need to look at each category separately. Starting with **profitability** we can separate:

#### Profitability Margins

These are essential for understanding how a company is managing its costs and generating profits from its revenues.

- **Net margin:** Represents net profit as a percentage of total sales, indicates a company's efficiency in generating profits after accounting for all expenses, taxes, and costs.
- **Operating margin:** Measures operating profit (EBIT) as a percentage of total sales, providing a clear view of the profitability of a company's core operations, excluding interest and taxes. This metric is fundamental for evaluating a company's operational efficiency.

- **EBITDA margin:** Removes the effects of capital structure and accounting policies, offering a clear view of the company's pure operational profitability.
- **Gross margin:** Focuses on revenues after deducting the cost of goods sold, is a key measure of production efficiency and a company's ability to manage its direct costs.

### Performance Ratios

These are used to measure the overall performance of the company.

- **ROA (Return on Assets):** Measures how efficiently a company converts its assets into profits. This is especially important in capital-intensive sectors, where efficient asset management can make a significant difference in profitability.
- **ROE (Return on Equity):** Focuses on the profits generated per dollar of equity invested by shareholders. This ratio is crucial for evaluating a company's overall profitability from the shareholders' perspective. It is an especially valuable metric for investors seeking to maximize their returns on equity investment.
- **ROIC (Return on Invested Capital):** Focuses on the return generated by all the funds invested in the company, including both shareholders' equity and debt. It is a comprehensive measure of a company's ability to generate value from all its sources of financing.
- **ROCE (Return on Capital Employed):** Measures the funds used to finance operations, regardless of the source. This ratio is useful for comparing the efficiency of companies with different capital structures, as it focuses on total capital employed rather than just equity.

Also, to not only look at actives, Tweenvest uses the cash generated to calculate: **CROIC** (Cash Return on Invested Capital), **CROCE** (Cash Return on Capital Employed), **OCF/Sales**, and **FCF/Sales**. And lastly it takes account also the Owner's income to calculate: **Owner's Income/Sales**, **Owner's CROIC** and **Owner's CROCE**.

Continuing with the **financial health**, we need to analyze the company's debt in different aspects:

### Leverage Ratios

These ratios assess how much a company relies on debt to finance its assets and operations, and are essential for evaluating financial risk and long-term solvency.

- **Financial Leverage** (Total Assets / Equity): Measures the proportion of a company's assets that are financed by shareholder equity. A higher ratio suggests the company is using more debt relative to equity, indicating greater financial risk but also potential return amplification through leverage.
- **Total Debt/Assets:** Indicates what portion of the company's assets is financed through debt. A lower ratio implies a more conservative capital structure, while a higher one may indicate increased risk if the company becomes over-leveraged.
- **Total Debt/Capital:** Measures the share of total capital (debt + equity) that comes from debt. This ratio is useful for understanding how dependent the company is on borrowed funds compared to its overall capital base.
- **Total Debt/Equity:** Compares the company's total debt to its shareholder equity. It provides insight into the balance between debt and equity financing. A high ratio may signal financial risk, but also the potential for higher returns if debt is managed well.

## Debt Coverage Ratios

These metrics evaluate a company's ability to cover its debt using its earnings or cash flow, reflecting the sustainability of a company's debt in relation to its operational performance.

- **Net Debt/EBIT:** Shows how many years it would take for a company to repay its net debt using EBIT (Earnings Before Interest and Taxes).
- **Net Debt/EBITDA:** Similar to the above, but adds back depreciation and amortization. This gives a more cash-focused view of a company's ability to handle its debt load, and is especially useful for comparing companies in capital-intensive industries.
- **Net Debt/FCF:** Evaluates how many years of free cash flow would be needed to pay off net debt. Since FCF includes investment needs, this ratio gives a more conservative view of debt sustainability.
- **Net Debt/Owner's Income:** Compares net debt to the income available to equity holders (after all operating and investing costs).

## Interest Coverage Ratios

These ratios measure how easily a company can meet its interest payments on outstanding debt — critical for assessing short-term debt service capability.

- **EBIT/Interest:** Indicates how many times a company can cover its interest expenses with its operating income.
- **EBITDA/Interest:** Similar to the above, but adds back depreciation and amortization. This gives a clearer picture of available cash earnings before fixed financial obligations, ideal for heavily asset-based businesses.
- **FCF/Interest:** Since FCF considers investment needs, this is a stringent test of how much real, discretionary cash is available for debt servicing.
- **Owner Earnings/Interest:** Evaluates a company's ability to meet interest payments based on the earnings effectively attributable to shareholders. It accounts for operational cash flow minus necessary capital expenditures.

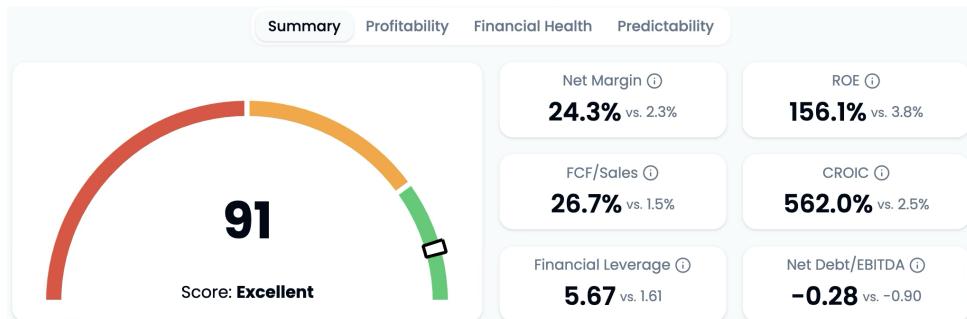
## Liquidity Ratios

These ratios measure a company's ability to meet short-term obligations with its short-term assets. They are essential for evaluating near-term financial health and risk of insolvency.

- **Current Ratio (Current Assets / Current Liabilities):** Shows whether a company has enough assets to cover its short-term liabilities. A value above 1 is generally considered healthy, though excessively high values may imply inefficiency.
- **Quick Ratio ((Current Assets - Inventory) / Current Liabilities):** A more stringent version of the current ratio that excludes inventory, which may not be easily liquidated. It's useful in assessing true short-term liquidity.
- **Cash Ratio (Cash and Equivalents / Current Liabilities):** The most conservative liquidity metric, focusing only on cash and equivalents. It shows the immediate solvency of a company in a worst-case scenario.
- **OCF Ratio (Operating Cash Flow / Current Liabilities):** Assesses how well the company's operational cash flows can cover its current obligations. This offers a realistic view of liquidity since it's based on actual cash generation rather than accounting figures.

## Predictability

And finally, for the quality score we need to look at the company's predictability. This is achieved by trying to fit values related to the company's success —such as Sales— to a exponential curve, using the ordinary least square method, which is supported by large financial literature [12].



**Figure 2.1:** Tweenvest's Quality Score

After calculating all of this ratios, Tweenvest compares them to the sectors' median and interpolates each of them to create a single score for the ratio, then it aggregates them all using personalized weights to create the final Quality Score, which is then showed to the clients as shown in Figure 2.1.

### 2.1.2 Growth

The Growth Score evaluates a company's historical growth across multiple key metrics, and compares them to industry standards. This comprehensive approach ensures a balanced assessment of growth across different aspects of the business:

#### Revenue & Profitability Growth

- **Sales:** Measures growth in core revenue streams
- **EBITDA:** Captures growth in operational cash-generating ability before non-cash and financing impacts
- **Operating Income:** Reflects growth in profit from core business operations
- **Net Income:** Includes all income sources, showing overall profitability growth

#### Cash Flow Growth

- **Operating Cash Flow:** Measures growth in cash generation from operations
- **Simple FCF:** A straightforward proxy for available cash after essential investments
- **Levered/Unlevered FCF:** Provide detailed views of free cash flow with and without debt impact
- **Owner Earnings:** Useful for volatile capex cases, emphasizing cash available to shareholders

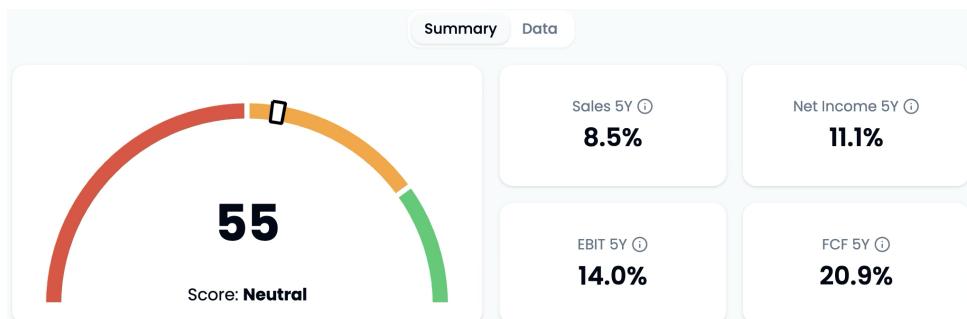
### Capital Base Expansion

- **Total Assets:** Indicates expansion in overall asset base
- **Equity:** Reflects growth in shareholders' claim on the business
- **Tangible Book Value:** Highlights growth in physical net assets, excluding intangibles
- **Invested Capital:** Captures total capital being put to productive use
- **Capital Employed:** A broader measure of capital supporting business operations

### Per-Share Value Growth

- **Diluted EPS:** Tracks per-share earnings growth, accounting for dilution effects
- **Diluted Shares:** Included to track share count changes, ensuring EPS growth isn't artificially inflated by buybacks or dilution
- **Ordinary DPS:** Tracks the growth of shareholder payouts, a proxy for confidence in future earnings

To compute the Growth Score, Tweenvest calculates 10-year, 5-year, and 3-year averages and then interpolates the growth rate to industry standards. This approach reinforces the long-term investment philosophy, giving lasting growing companies a better score.



**Figure 2.2:** Tweenvest's Growth Score

### 2.1.3 Value

The Value Score measures how attractively a company is priced relative to fundamental multiples such as earnings, cash flow, sales, and dividends. This is critical for investors following value investing principles, where the goal is to buy quality companies for less than their intrinsic worth. The algorithm evaluates a series of valuation multiples—both price-based and enterprise value-based—and dividend yield.

#### Price-Based Multiples

- **P/E (Market Cap / Adjusted TTM Earnings):** Measures how much investors are willing to pay per dollar of earnings
- **P/S (Market Cap / TTM Revenue):** Useful when earnings are volatile; shows valuation relative to sales
- **P/CF (Market Cap / TTM Operating Cash Flow):** Reflects valuation relative to cash-generating ability

- **P/B** (Market Cap / Tangible Equity): Especially relevant for asset-heavy sectors like banks or industrials

### Enterprise Value-Based Multiples

- **EV/Sales** (Enterprise Value / TTM Revenue)
- **EV/EBITDA** (Enterprise Value / TTM EBITDA)
- **EV/EBIT** (Enterprise Value / TTM EBIT)
- **EV/FCF** (Enterprise Value / TTM Free Cash Flow)

### Yield-Based Valuation

- **Dividend Yield (%)** (Dividend per Share / Price per Share)

Each of these ratios is compared to multiple historical statistics and sector benchmarks to create individual scores and then average them.

#### Note:

- $\text{Market Cap} = \text{Share Price} \times \text{Total Outstanding Shares}$
- $\text{Enterprise Value} = \text{Market Capitalization} + \text{Total Debt} - \text{Cash} + \text{Marketable Securities}$

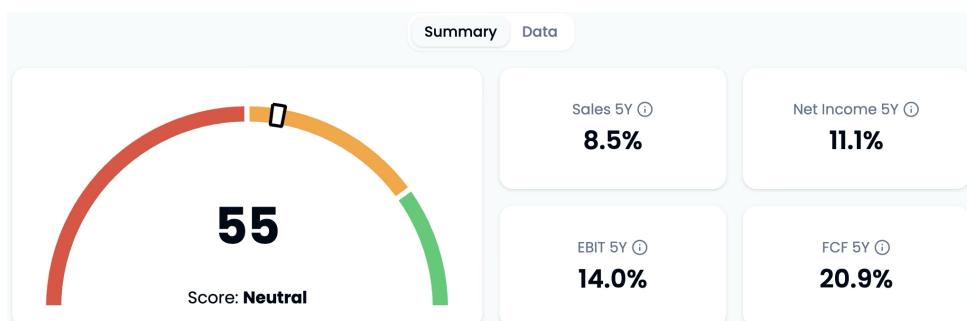


Figure 2.3: Tweenvest's Value Score

#### 2.1.4 Dividend

When analyzing Tweenvest most common investors profile, we see a high tendency to income-focused and long-term investors using mainly dividends as their principal concern. Thats why the platform dedicated a full section to this topic.

The Dividend Score measures the attractiveness, reliability, and growth potential of a company's dividend payments. It helps investors assess whether the dividend is both rewarding today and sustainable for tomorrow. To do so, the score was built from three primary components:

## Safety

This part of the score is used to assess the sustainability of the dividend.

- **Payout Ratio EPS** (DPS / Diluted EPS): Shows if dividends are covered by accounting earnings
- **Payout Ratio FCF** (DPS / Free Cash Flow per Share): Shows if dividends are funded by real cash generation
- **Payout Ratio Owner Earnings** (DPS / Owner Earnings per Share): A conservative test of sustainability (excludes CAPEX)

## Growth

Measures how consistently and strongly the dividend has grown over time.

- **Ordinary DPS CAGR**: 3-Year, 5-Year, 10-Year growth rates

## Yield

This evaluates the attractiveness of the dividend today relative to the company's historical averages and sector benchmarks.

- **Dividend Yield** (DPS / Price per Share): Represents how much income an investor receives annually from dividends

These components are individually scored weighted and interpolated against industry benchmarks to form a composite score.



Figure 2.4: Tweenvest's Dividend Score

Finally, as seen in Figure 2.4, the algorithm adjusts this score based on how many years the dividend has been maintained or increased, **rewarding consistency**.

## 2.2 Code practices

Since the code that needs to be changed is used in a production environment by Tweenvest, it is important to follow some *good practices* to ensure the code is easy to understand, maintain, and to avoid introducing new bugs.

- **Understanding and Documenting**: Before starting to work on the code, it is important to understand the codebase and the purpose of the code. For this, it is recommended to use some tools like flux diagrams, code comments, and documentation. As it will be shown later on.

- **Testing:** To ensure proper functionality, every function and class should have unit tests that cover all possible scenarios. These tests are run automatically by the CI/CD pipeline.
- **Pair code review:** After completing the previous steps, the code undergoes a pair code review process, where at least one other team member reviews and approves the changes in *Github* before merging into the main branch.
- **Logging:** After each feature is implemented, it is important to analyze the actual performance in the production environment by looking at its logs to check if the feature is working as expected, and to look for any possible optimizations to be made.

## 2.3 Preprocessing

Once we have created our raw dataset, we need to inspect it and make sure it is consistent and ready to be used for the modeling. This part is really important to avoid any bias in the modeling process and to only use the data that is actually relevant for the analysis.

### 2.3.1 Data Consistency

Since all of the data is coming from a data provider that uses OCR as one of their tools whenever they don't have the data in a structured format, we need to be very restrictive with the data used. Also, as we mentioned before, some of the algorithms for calculating the scores use long-term data such as 10 year growths rates.

This is one of the biggest **limitations to the study**: we will be missing the data of companies that stopped existing or that were acquired by other companies, if they didn't exist for 10 or more years.

For inspecting the data, we will use the *pandas* [15], *matplotlib* [9], and *seaborn* [16], which are powerful tools for data manipulation and analysis. The focus will be on the following aspects:

- **Missing Values:** We need to check if there are any missing values in the data.
- **Data Distributions:** We need to make sure that the data is distributed in a way that is suitable for the modeling.
- **Variable Correlations:** We also have to see if there are any correlations between the variables. For this, we will use the correlation matrix defined as:

$$\rho_{X,Y} = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y} \quad (2.1)$$

where  $\text{Cov}(X, Y)$  is the covariance between the variables  $X$  and  $Y$ ,  $\sigma_X$  is the standard deviation of  $X$ , and  $\sigma_Y$  is the standard deviation of  $Y$ . This matrix will be created with a custom heatmap function (see Appendix 7.4).

- **Data Patterns:** We need to check if there are any visible data patterns between the variables. To address this, we will use a custom *pairplot* function from *seaborn*, which is a powerful tool for data visualization and exploration depending on the type of variable (see Appendix 7.4).

Additionally, we decided to assign a value of 0 to all null Dividend scores, as a null value indicates that the company does not pay dividends. This choice was made to ensure that the Dividend score remains part of the analysis when evaluating its relevance to long-term profitability, rather than being excluded due to missing values.

### 2.3.2 Data Transformation

To ensure a robust and well-performing model, it is often necessary to transform the data into a format more suitable for learning algorithms, with the obligatory condition of being able to revert it back to its original scale when needed.

This step is particularly important because many components of machine learning models—such as the RBF kernel in Support Vector Machines or the regularization terms (L1 and L2) in linear models—implicitly assume that features are centered around zero and have comparable variances. Without this adjustment, features with significantly larger variances could dominate the optimization process, leading the model to underweight or ignore more informative but lower-variance features.

#### Standard Scaler

When training the neural networks, it is really important to standardize the data. We can easily use the *StandardScaler* from the *scikit-learn* library that transforms a feature  $x_i$  with the following formula:

$$z_i = \frac{x_i - \mu}{\sigma} \quad (2.2)$$

where  $\mu$  is the mean of the feature and  $\sigma$  is its standard deviation. This transformation results in a distribution with a mean of 0 and a standard deviation of 1.

#### Gaussian Transformation

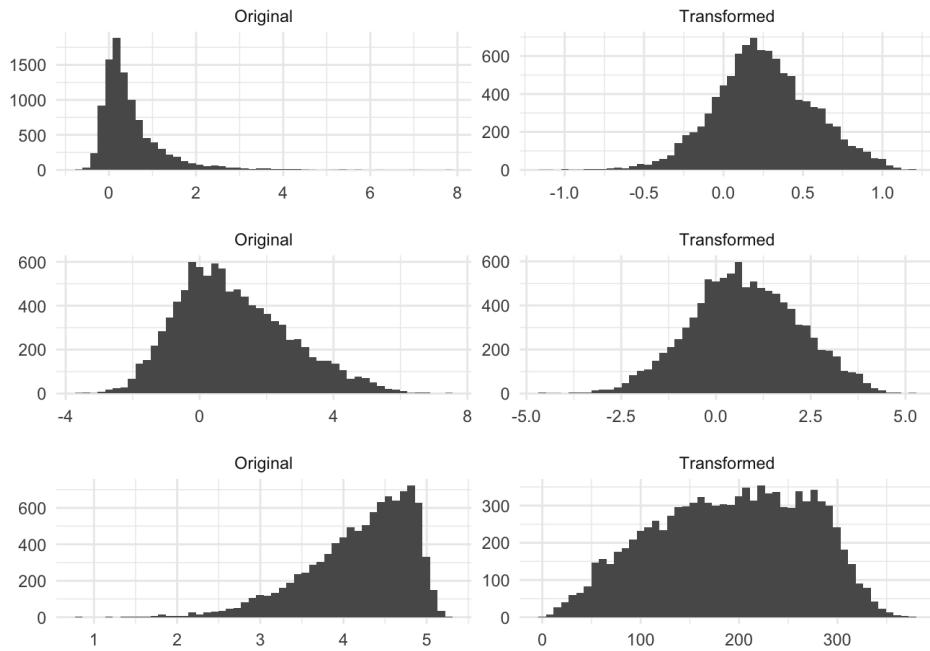
In numerous modeling applications, having normally distributed features is advantageous. **Power transformations** represent a set of parametric, monotonic functions that are an extension of the Box-Cox transformation. They are designed to convert data from various distributions into approximately Gaussian distributions, thereby reducing variance fluctuations and decreasing distribution asymmetry.

We decided to use the **Yeo-Johnson transformation** because it allows for negative values and it is reversible.

$$x_i^{(\lambda)} = \begin{cases} \frac{[(x_i+1)^\lambda - 1]}{\lambda}, & \text{if } x_i \geq 0, \lambda \neq 0 \\ \ln(x_i + 1), & \text{if } x_i \geq 0, \lambda = 0 \\ -\frac{[(-x_i+1)^{2-\lambda} - 1]}{2-\lambda}, & \text{if } x_i < 0, \lambda \neq 2 \\ -\ln(-x_i + 1), & \text{if } x_i < 0, \lambda = 2 \end{cases} \quad (2.3)$$

Where  $\lambda$  is a power parameter that helps minimize the skewness of the data. And since we have already deleted extreme outliers, the final distribution shouldn't be too distorted.

Here are some examples of the transformations for different distributions:



**Figure 2.5:** Yeo-Johnson Transformations examples

## 2.4 Outlier Detection

Anomalies are data patterns that have different data characteristics from normal instances. The ability to detect anomalies has significant relevance since often provide critical and actionable information in many different contexts.

For example, anomalies in credit card transactions could signify fraudulent use of credit cards, or an unusual computer network traffic pattern could stand for an unauthorised access.

### 2.4.1 Inter Quartile Range

Given a univariate dataset, the Interquartile Rule identifies outliers based on the interquartile range (IQR). The steps are as follows:

1. Compute the first quartile ( $Q_1$ ), which is the 25th percentile of the data.
2. Compute the third quartile ( $Q_3$ ), which is the 75th percentile of the data.
3. Calculate the interquartile range:

$$\text{IQR} = Q_3 - Q_1$$

4. Define the lower and upper bounds for non-outlier values following the 1.5 rule:

$$\text{Lower bound} = Q_1 - 1.5 \times \text{IQR}$$

$$\text{Upper bound} = Q_3 + 1.5 \times \text{IQR}$$

5. Any data point  $x$  is considered an outlier if:

$$x < \text{Lower bound} \quad \text{or} \quad x > \text{Upper bound}$$

As we can see this is a very simple and intuitive method, but it has some limitations:

Pros	Cons
Simple to compute and interpret.	Not suitable for multivariate data with correlated features.
Non-parametric.	May misclassify values in skewed distributions as outliers.
Robust to extreme values, since it relies on percentiles rather than the mean.	Fixed multiplier is arbitrary and may need tuning for each dataset.

**Table 2.1:** Pros and Cons of the Interquartile Rule

Even though this method is simple to compute and interpret, it is not suitable for multivariate data with correlated features, and that's where the rest of the methods come into play.

#### 2.4.2 Single Vector Machine

Support Vector Machines (SVMs) represent a robust class of unsupervised learning algorithms that can be effectively adapted for anomaly detection tasks. This algorithm establishes itself on the premise that the majority of real-world data is inherently normal. Where the goal is to define a boundary encapsulating the normal instances in the feature space, thereby creating a region of familiarity.

To capture the bases of the *scikit-learn* implementation [4] used in this work, we can say that the main idea is to find a minimum volume sphere that contain all the training samples. This sphere, described by its center  $c$  and its radius  $r$ , is obtained by solving the constrained optimization problem: [17]

$$\min_{r,c} r^2 \quad \text{subject to} \quad \|\Phi(x_i) - c\|^2 \leq r^2 \quad \text{for } i = 1, 2, \dots, n \quad (2.4)$$

This boundary is strategically positioned to maximize the margin around the normal data points, allowing for a clear delineation between what is considered ordinary and what may be deemed unusual. This emphasis on margin maximization is akin to creating a safety buffer around the normal instances, fortifying the model against the influence of potential outliers or anomalies.

In summary, this method has the following characteristics:

Pros	Cons
Effective for high-dimensional data and complex boundaries.	Sensitive to the choice of kernel and hyperparameters (e.g., $\nu$ , $\gamma$ ).
Works well when the training set contains mostly or only normal instances.	Computationally intensive, especially on large datasets.
Can capture nonlinear patterns using kernels (e.g., RBF kernel).	Difficult to interpret or explain the decision function.
No need for labeled data from the outlier class.	-
Solid theoretical foundation and widely supported in machine learning libraries.	-

**Table 2.2:** Pros and Cons of One-Class SVM for Outlier Detection

### 2.4.3 Isolation Forest

As we are seeing in this small fraction of outlier detection methodologies, most of the existing anomaly detection approaches are based on the premise of normal distributions, then identify anomalies as those that do not conform to the normal profile.

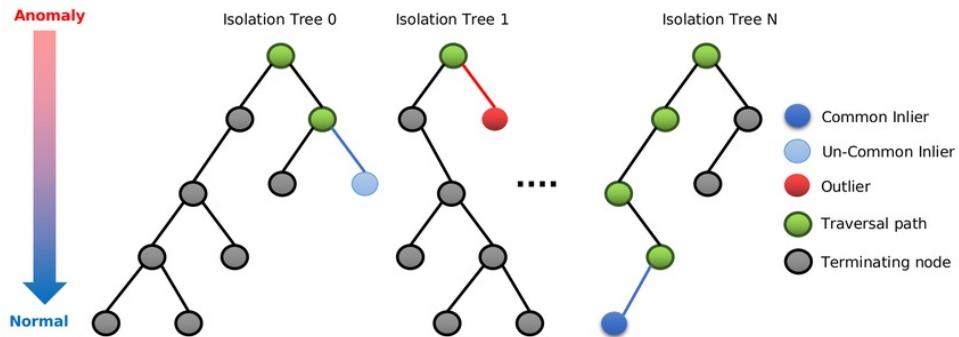
But to solve this issue we can use the Isolation Forest algorithm, which is explained in the original paper Liu, Ting, and Zhou [10]. The algorithm constructs multiple isolation binarytrees (iTrees) from the dataset, where anomalies are identified as data points that exhibit shorter average path lengths across these trees.

Lets consider a dataset  $X = \{x_1, \dots, x_n\}$  containing  $n$  points in  $d$ -dimensional space, and a subset  $X' \subset X$ . An Isolation Tree (iTree) is constructed as follows:

- Each node  $T$  in the tree is either:
  - An external node (leaf) with no children, or
  - An internal node with exactly two children ( $T_l$  and  $T_r$ ) and a test condition
- The test condition at each internal node consists of:
  - A randomly selected attribute  $q$
  - A split value  $p$
  - A test  $q < p$  that determines whether a point goes to  $T_l$  or  $T_r$

The tree construction process recursively partitions  $X'$  by randomly selecting attributes and split values until either:

- The node contains only one instance, or
- All instances at the node have identical values



**Figure 2.6:** Visualization of how Isolation Forest works.

As we can see in Figure 2.6, once the iTree is fully constructed, each point  $x_i \in X$  is isolated at a leaf node. The path length  $h(x_i)$  of a point  $x_i$  is defined as the number of edges traversed from the root to its leaf node. Points with shorter path lengths are considered more likely to be anomalies, as they require fewer splits to be isolated from the rest of the data.

This method works very good in most cases, but it has also its own limitations:

Pros	Cons
Efficient and scalable to large datasets due to its tree-based structure.	Less effective for detecting local outliers in densely clustered regions.
Non-parametric: does not assume any data distribution.	Performance can vary with the choice of parameters like number of estimators and subsample size.
Naturally handles high-dimensional data.	Not suitable for detecting subtle anomalies that closely resemble normal data.
Requires little preprocessing (no need for feature scaling).	–
Robust to irrelevant features due to random sub-sampling.	–

**Table 2.3:** Pros and Cons of the Isolation Forest Method for Outlier Detection

#### 2.4.4 Local Outlier Factor

Since we are looking at outliers from the perspective of correlations between the different response variables to see which companies don't behave naturally. We also checked on the Local Outlier Factor (LOF) for being based on hyper-plane densities of data.

The algorithm measures the local deviation of the density of a given sample with respect to its neighbors. It is local in that the anomaly score depends on how isolated the object is with respect to the surrounding neighborhood. More precisely, locality is given by k-nearest neighbors, whose distance is used to estimate the local density. By comparing the local density of a sample to the local densities of its neighbors, one can identify samples that have a substantially lower density than their neighbors. These are considered outliers. Breunig et al. [2]

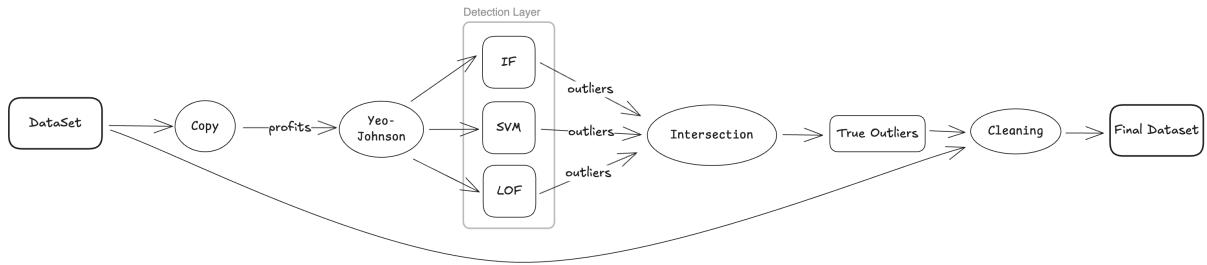
Pros	Cons
Can detect outliers relative to their local neighborhood, allowing it to identify context-specific anomalies.	LOF scores are relative and lack a universal interpretation threshold for outliers.
Effective in datasets with varying densities — unlike global methods, it can recognize outliers near dense clusters.	Sensitivity to parameter selection (e.g., number of neighbors) can affect performance and consistency.
Applicable in any domain where a dissimilarity measure is defined, not limited to vector spaces.	Not inherently scalable to very large datasets without approximation or optimization.
Works well across domains, such as network intrusion detection or classification tasks.	–
Easily generalizable and adaptable for use in spatial data, temporal data, and network structures.	–

**Table 2.4:** Pros and Cons of the Local Outlier Factor (LOF) Method

#### 2.4.5 Multi-Criteria Outlier Detection

There is a big controversy in the field about the best method to detect the outliers, since the results can vary a lot depending on the algorithm used.

So to attempt to fix this issue, after reading some aggregative methods such as Abro, Taşci, and Uğur [1], we created a multi-criteria method that combines the results of the different methods to get a more robust result.



**Figure 2.7:** Multi-Criteria Outlier Detection Method

As shown in the Figure 2.7, this starts by homogenizing the variables distribution with a common data transformation so they are more Normal, but to make sure that the final data isn't being transformed we copy the original dataset earlier. After this, we use multiple methods to detect the different outliers and then we make the intersection between them to get the common outliers. By the end of this process we have tagged the original dataset rows, assuring that the data that will be excluded is only formed by "*true outliers*".

We noticed that since this method is very restrictive, we have to set a threshold multiplier to the original dataset to get the final outliers. So for example, a 20% threshold would only exclude around 4% of the original dataset because of the intersection of all the methods makes it very restrictive. Here we **propose for future work** to fine tune the different thresholds for each method to get a more robust result. This will be explained with more detail in the development section.

## 2.5 Predictive Models

The main objective of this thesis is to check if the scores have any relation with the future profits of the companies. For doing this in the most generalized way, instead of creating a portfolio following some subjective criteria on the scores, markets and other variables, we used predictive modeling to check for the actual relationships between the response variables and the explanatory ones.

As a first approach, we created a set of models that separate the scores and the different returns.

### 2.5.1 Regression Models

To create the econometrical models, we want to be able to understand the actual relationship between the response variables and the explanatory ones. For this we started with the creation of multiple regression models.

#### Linear Regression

As for the first models, we created a series of linear regression models for each factor score that includes the interactions between the dummy variables and the score used in each model.

$$\hat{P}_i = \beta_0 + \beta_1 \cdot S_j + \sum_{n=2}^N \beta_n \cdot D_n + S_j \cdot \sum_{m=N+1}^M \beta_m \cdot D_m \quad (2.5)$$

where:

- $P_i$  is the estimated profit for each time horizon  $i$
- $S_j$  is the score for each factor  $j$
- $D_{n,m}$  are the dummy variables for the industry and region, where  $D_{n,m} := I_k \cup R_l$
- $\beta$  are the coefficients that weight each variable.

For fitting the models, we also applied a backward selection to remove the variables that are not statistically significant, using the **p-value** criterion.

### Generalized Additive Models

Because the relationships between the variables are not linear, we also used the Generalized Additive Models (GAM) to check if they were able to capture the non-linear relationships between the variables.

According to Servén and Brummitt [13], GAMs are smooth semi-parametric models that can capture non-linear relationships between variables. They take the form:

$$g(\mathbb{E}[y|X]) = \beta_0 + f_1(X_1) + f_2(X_2, X_3) + \cdots + f_M(X_N) \quad (2.6)$$

where:

- $X^T = [X_1, X_2, \dots, X_N]$  are the independent variables
- $y$  is the dependent variable.
- $g()$  is the link function that relates the predictor variables to the expected value of the dependent variable.
- $f_m()$  are feature functions built using penalized B-splines, which automatically model non-linear relationships without requiring manual transformation of variables.

In our case, we used pyGAMs LinearGAM since the model gives a Normal error distribution, and an identity link.

### 2.5.2 Time Series

In the stock market, the prices of the stocks are not independent of each other, they are correlated in the time series and they also have memory on past behavior of the company. This is what is known as **Momentum**, so companies that have a good past performance are more likely to have a good future performance, and vice versa.

For this reason, we contemplated two possibilities:

#### ARIMA Models

To take account the profit tendency, we implemented ARIMA models to improve the regression models performance by using the residues of the already fitted model:

$$P_i = \hat{P}_i + \varepsilon_i \quad \longleftrightarrow \quad \varepsilon_i \sim ARIMA_i(p, d, q) \quad (2.7)$$

The ARIMA model is defined as:

$$\phi_i(B_i)(1 - B_i)^{d_i} \varepsilon_i^t = \theta_i(B_i) \eta_i^t \quad (2.8)$$

where for each  $i$  profit:

- $\varepsilon_i^t$  is the residual at time  $t$ .
- $\eta_i^t$  is a white noise error term (innovation).
- $B_i$  is the backshift operator, such that  $B\varepsilon_i^t = \varepsilon_i^{t-1}$ .
- $\phi_i(B) = 1 - \phi_{i1}B - \phi_{i2}B^2 - \dots - \phi_{ip}B^p$  is the autoregressive (AR) polynomial.
- $\theta_i(B) = 1 + \theta_{i1}B + \theta_{i2}B^2 + \dots + \theta_{iq}B^q$  is the moving average (MA) polynomial.
- $d_i$  is the order of differencing.

### Windowed Models

And for capturing possible memory of the companies behavior in different aspects, we also implemented windowed models that used the last  $N$  scores statistical properties to predict future profits.

So for each score  $j$ , we will calculate the average and the standard deviation of the scores over the last  $N$  periods. Let  $S_j^t$  be the score at time  $t$ , then for a window of size  $N$ :

$$\bar{S}_j^N = \frac{1}{N} \sum_{i=t-N+1}^t S_j^i \quad (2.9)$$

where:

- $\bar{S}_j^N$  is the average score over window  $N$  for score  $j$
- $N$  can be 3 months, 6 months, 1 year, or 2 years.
- $t$  is the current time point

For adding also the deviation from the average, we also calculated the standard deviation of the scores over the same window:

$$\sigma_j^N = \sqrt{\frac{1}{N} \sum_{i=t-N+1}^t (S_j^i - \bar{S}_j^N)^2} \quad (2.10)$$

These statistical measures will be used as additional features in our predictive models to capture the temporal behavior of each score.

### 2.5.3 Neural Networks

Finally, for trying to capture the non-linear relationships between all the available variables, including the windowed statistical properties, we will use the Neural Networks.

## 2.6 Back-testing

In order to assess and compare the modeling techniques employed in this thesis, it is important to have a consistent criteria to evaluate them based on the next requirements:

## 2.6.1 Performance

The first thing we need to see is if the models are able to detect any relationships between the co-variables and the response variable, and how accurate they are:

- **Prediction Accuracy:**  $R^2$  and  $R^2_{adjusted}$ , these metrics quantify the proportion of variance explained by the model.
- **Accuracy:** RMSE, MAE, which evaluate the errors the model makes when predicting the target variable.

## 2.6.2 Consistency and Overfitting

We check whether the models reflect stable, interpretable patterns in the relationships between scores and alpha. For this, we will use the following techniques:

- **Sign of coefficients:** We will check if the sign of the coefficients remains stable across different time horizons.
- **Shape of functions:** We will check if the shape of the functions remains stable across different time horizons.
- **Number of significant variables:** We will check if the number of significant variables increases with the performance of the models, using the  $p-value < 0.2$  criterion.

## 2.6.3 Robustness of Results

The residuals of the models should resemble white noise, indicating the absence of unmodeled structure and valid statistical inference Enders [8].

- **Mean Value Test:** The mean of the residuals should be close to 0, which is checked with the T-test.
  - **Distribution Plots:** Visualize the distribution of residuals to assess if the mean is centered at zero.
- **Heteroscedasticity Tests:** The variance of the residuals should be constant across all time horizons.
  - **White Test:** Checks for general forms of heteroscedasticity.
  - **Breusch-Pagan Test:** Tests whether the variance of the errors depends on the values of the independent variables.
  - **F and t tests:** Used to compare variances and means across groups.
  - **Distribution Plots:** Visualize residuals versus fitted values to detect patterns in variance.
- **Normality Tests:** The distribution of residuals should follow a Normal distribution.
  - **Shapiro-Wilk Test:** Statistical test for normality.
  - **Jarque-Bera Test:** Tests whether sample data have the skewness and kurtosis matching a normal distribution.
  - **Q-Q Plots:** Visual tool to compare the quantiles of the residuals to a normal distribution.
  - **Distribution Plots:** Visualize the histogram or density of residuals.

- **Autocorrelation Test:** The errors should not be correlated with each other.
  - **Durbin-Watson Test:** Checks for the presence of autocorrelation in the residuals.
  - **ACF and PACF Plots:** Visualize the autocorrelation and partial autocorrelation of residuals to detect serial correlation.

# Chapter 3

## Development

### 3.1 Architecture Enhancement

At the current moment of beginning this work, Tweenvest kept a large DB of most of the needed information for creating the dataset. This included all of the price histories, historical currency multipliers to dollar, dividends payed per stock... but when it came to the scores we had a traceability problem.

For optimizing the costs and structure of the DB, Tweenvest chose to only save the latest score for each company and overwriting it each day when calculating the newest one, which led to the first main task.

#### 3.1.1 Updating Tweenvest Code

After locating where the score calculators are called, we have to understand the code's architecture and the relationships between the different database models to see exactly what needs to be changed. Tweenvest uses Django as the main backend's technology to handle a relational database with PostgreSQL, this is due to the nature of the financial data where one the information is related to another.

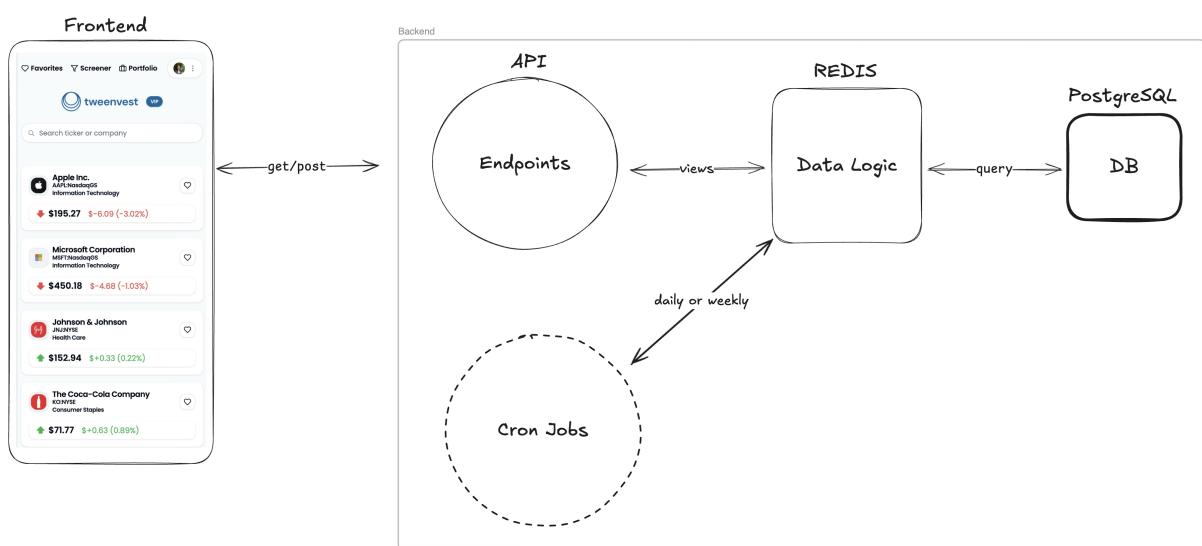
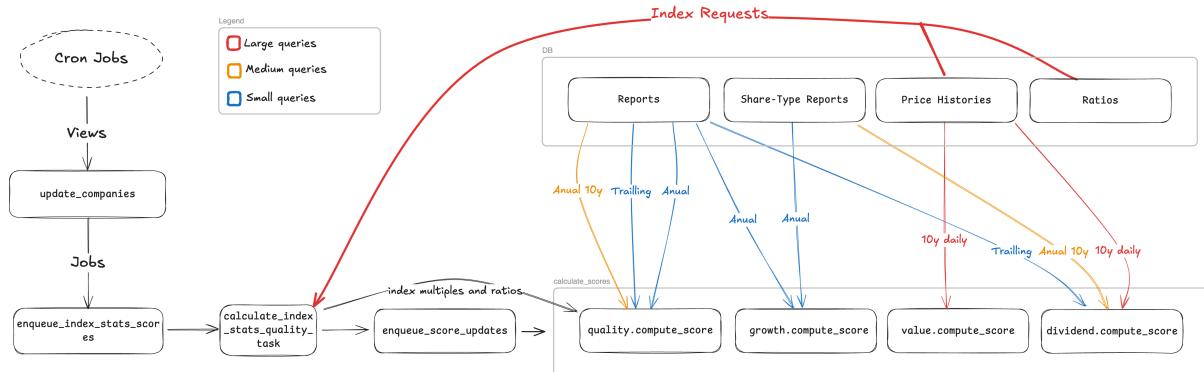


Figure 3.1: General Tweenvest's Architecture

As we can see, there are two different actions occurring at the same time. For one side we have all of the requests coming from the users interactions, and for the other hand the programmed

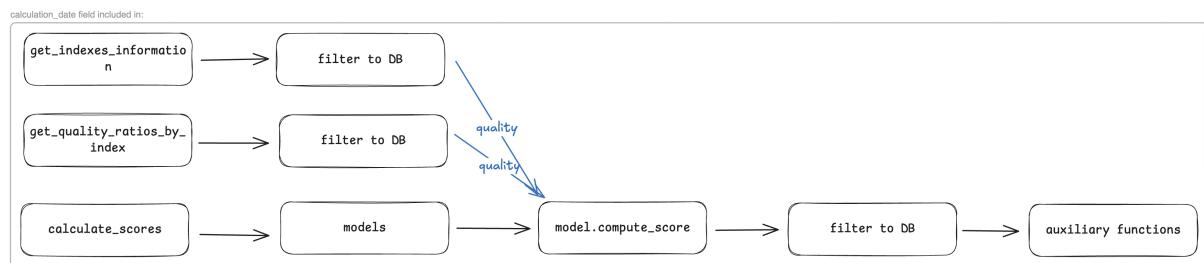
jobs that need to be executed every day to update all of the financial information of the stocks, or necessary actions such as sending emails to get authentication pins.

Since we are changing the data logic for calculating and storing historical data of the factor scores, we need to assure that the systems capabilities don't exceed the platforms needs, and that the changes don't affect the normal calculation of the factor scores, so we had to look at the code corresponding to the factor scores.



**Figure 3.2:** Scores Calculations Code Schema

With this schema we have simplified the logic behind several functions related to the calculations of each factor score. Quickly we realized that we had to modify the queries to include an extra filter, *calculation\_date* to get the models related to the score at any time, this way we can set it as "today" whenever we are calculating the present scores.



**Figure 3.3:** *calculation\_date* Field Propagation Schema

Following the propagation schema in Figure 3.3, the root change comes from the functions *calculate\_scores*, *get\_indexes\_information* and *get\_quality\_ratios\_by\_index*. These are the ones used by the production cron jobs to calculate the factor scores daily, so for a normal function call we set the *calculation\_date* to None and whenever we want to calculate the scores for a specific date we set it to the desired date.

This helped us create very readable logic to adapt the filters to both scenarios, for example in the value score shown in Listing 3.1:

```

1 if calculation_date is None:
2     calculation_date = date.today()
3     filters = {
4         "stock": stock,
5         "date__gte": calculation_date - relativedelta(years=10),
6     }

```

```

7 else:
8     filters = {
9         "stock": stock,
10        "date__gte": calculation_date - relativedelta(years=10),
11        "date__lte": calculation_date,
12    }
13 phs = (
14     PriceHistory.objects.filter(**filters)
15 )

```

**Listing 3.1:** Value Score Filtering

So as we can observe, the main change for this part of the code was to add a date filter to get only data "lower than or equal" to the *calculation\_date*. This may seem a bad approximation, but it has to be done this way due to the nature of the data: most of the financial data come from annual, semestral, and quarterly reports, except to the price history that is updated daily.

Along the way of analyzing the code's structure we made some adjustments to simplify the legibility:

- **Renaming functions** whenever they are internal methods or generic auxiliary methods called by different functions.
- **Externalizing mocking functions** that help create fake models and data for unit testing.
- **Homogenize the code's structure**, field types, to help with readability and keeping a standard order along the backend and frontend.

Another very important part of these updating procedure was to **adapt the unit tests** to make sure that the original code still behaved as expected, and creating new ones to check that the changes calculated the factor scores properly.

### 3.1.2 Designing new Jobs

Once all tests have passed and the team had approved the changes, the next step was to create scripts —from now on, *jobs*—to populate the current DB with the historic factor scores while maintaining the servers performance in normal levels. For this, we had to update a basic method, *bulk\_create\_or\_update*, used many times through the whole codebase and add a new conditioning so it can handle conflicts depending if the models are empty or not. After reading the official Django Documentation [5], we saw that in newest versions they had included new fields to its original method *bulk\_create*, so we had to update the Django version in our Docker machine and make sure this didn't introduce any unwanted issues.

After this was implemented, we decided to create three different jobs for enqueueing the tasks and being able to separate the calculations. This way we could restart them in case something broke during the process (see Appendix 7.1). Also, we had to manage the server resources, so for assuring the platform's performance, we used the following hierarchy (see Listing ??):

1. Seven workers for primary daily jobs from different API data integration
2. One dedicated worker for guaranteeing email sendings, and 6 prioritized shared workers
3. Four shared workers with lower priority, except for 1 which is set higher than email.

A crucial part of this code was replicability, equal distribution through different sectors, and mostly to assure that the companies existed during enough time for the factor scores to calculate the 10 years averages.

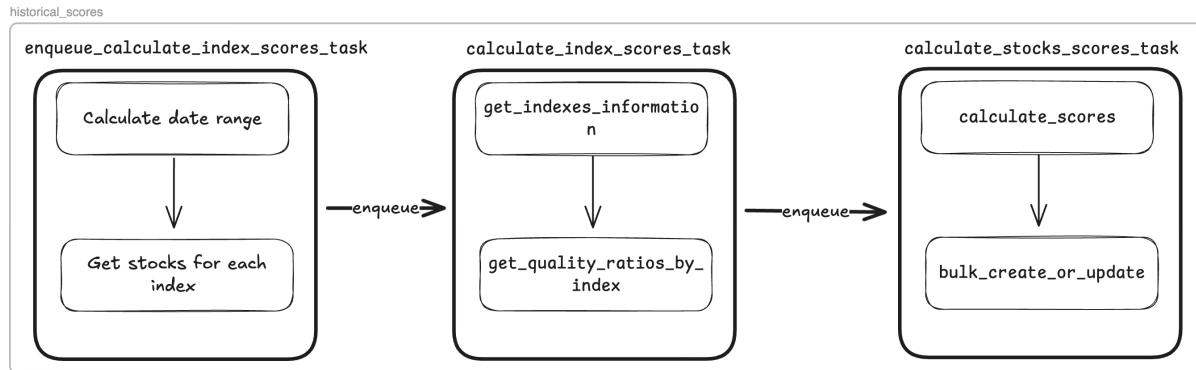


Figure 3.4: Historical Jobs First Iteration Schema

We created the first jobs with the schema shown in Figure 3.4, and during the first small test it worked in local, but when tested with the production DB we started to see important issues with how the functions where designed due to the large amount of data being processed:

- Passing incorrect arguments between jobs—such as entire lists of stock objects—led to the REDIS database reaching its maximum storage capacity.
- Jobs ran out of time.

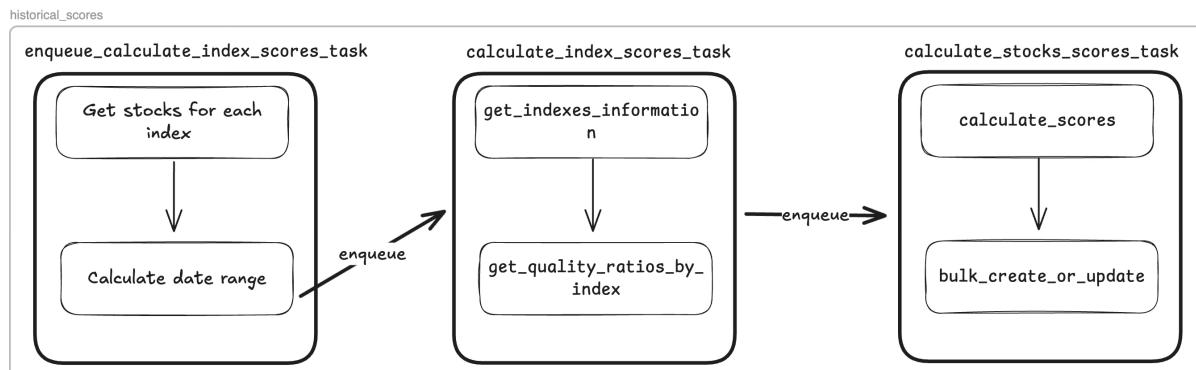


Figure 3.5: Historical Jobs Final Version Schema

We then changed the day range logic to enqueue a *calculate\_index\_scores\_task* for each day, so the job didn't have to calculate all of the index data for the dates at all at once, as shown in Figure 3.5. Also, we modified the function args to only uses models identifiers "id" to reduce the REDIS memory used. The complete implementation of this job can be found in Appendix 7.1.

### 3.1.3 Telemetry Tracing

Finally we had the functional jobs for calculating the scores, and ran the first calculation for a dataset of 1,500 stocks. After the first jobs finished we quickly realized that they were taking too much time to process, and saw a disproportionate amount of queries to the DB so we started a process to check on where was the bottleneck using telemetry tracing tools.

To do so, we had to inspect each method and function that was being used through the calculations of the factor scores. Luckily we had just recently implemented *DataDog* which is a software that provides monitoring and analytics for applications and infrastructure, offering real-time metrics, event monitoring, and end-to-end tracing.

To apply the tracer, we simply had to add a specific decorator before each method and set up the environment for being able to trace not only production launched jobs, but also development tests.

Here's an example of how the tracer was implemented in the quality score calculation:

```

1 from opentelemetry import trace
2
3 tracer = trace.get_tracer(__name__)
4
5 @tracer.start_as_current_span("compute_quality_score")
6 def compute_quality_score(
7     stock: "Stock",
8     index_quality_ratios: Dict[str, IndexRatio],
9     calculation_date: date | None,
10 ) -> tuple[dict, dict]:
11     ...
12     return final_score

```

**Listing 3.2:** Telemetry Tracing Implementation

We then looked at the *DataDog* dashboard and clearly saw an excess of queries to the DB when calculating each score, and tried to reduce them by looking for a simple fault in the logic of the queries. Similar issues had been resolved before by adding a *select\_related* in the queries. To clarify what this does, here is the definition with a simple use-case:

"Returns a QuerySet that will follow foreign-key relationships, selecting additional related-object data when it executes its query. This is a performance booster which results in a single more complex query but means later use of foreign-key relationships won't require database queries." – Documentation [6].

In Django, *select\_related()* and *prefetch\_related()* are designed to stop the deluge of database queries that are caused by accessing related objects. *select\_related()* "follows" foreign-key relationships, selecting additional related-object data when it executes its query. *prefetch\_related()* does a separate lookup for each relationship and does the "joining" in Python. So one uses *select\_related* when the object that you're going to be selecting is a single object, so *OneToOneField* or a *ForeignKey*.

After numerous attempts to improve performance, we were forced to reconsider our approach and began analyzing the problem from a different perspective. It quickly became evident that we were operating on a small virtual machine (Table 3.1) shared with the regular platform workload. Tweenvest's servers are not designed for high computational demand; rather, they are optimized for routine operations, where standard score calculations typically take between 3 to 4 hours. This setup works well under normal conditions, as calculations are performed after market hours and do not impact the user experience. However, it proved inadequate for our needs—specifically, processing data over five or more years within a reasonable timeframe, even when computing just one score per month.

VCPUS	RAM	SSD Storage	Extra Volume
8	4 GB	320 GB	–

Table 3.1: Tweenvest Server Specifications

### 3.1.4 Dedicated Server Deployment

We made the decision to establish a dedicated server for this thesis, equipped with enhanced computational power to facilitate the calculations and data fitting. After evaluating several providers, we opted for Hetzner servers due to their exceptional quality-price ratio. The final setup consists of a dedicated server with the following specifications:

VCPUS	RAM	SSD Storage	Extra Volume
16	32 GB	320 GB	480 GB

Table 3.2: Hetzner Server Specifications

Deploying the project in a scalable and reproducible environment involved several structured steps, encompassing server provisioning, secure remote access configuration using SSH, and initializing the containerized application deployment via Docker, with the application code-base managed through GitHub.

After all of the setup was done, we saw an amazing **x100 performance improvement**, going from 15-minute calculation times for index statistics to 8 seconds. Finally, we could launch the factor scores historical calculations, so we created two subsets of 1,500 random companies each with the factor scores calculated for the 1st day of the month for the following periods: 2015-2020 and 2020-2025, which took a total of 3 hours.

Now that we had filled the DB with the necessary factor scores for the study, we ran a big calculation to get all of the possible factor scores for the +100,000 companies, which took several days.

## 3.2 Dataset Creation

### 3.2.1 Aggregating the data

For being able to continue, we needed to create the final dataset with all the necessary data for later on doing the analysis and developing the predictive models, so we created a simple algorithm with this logic:

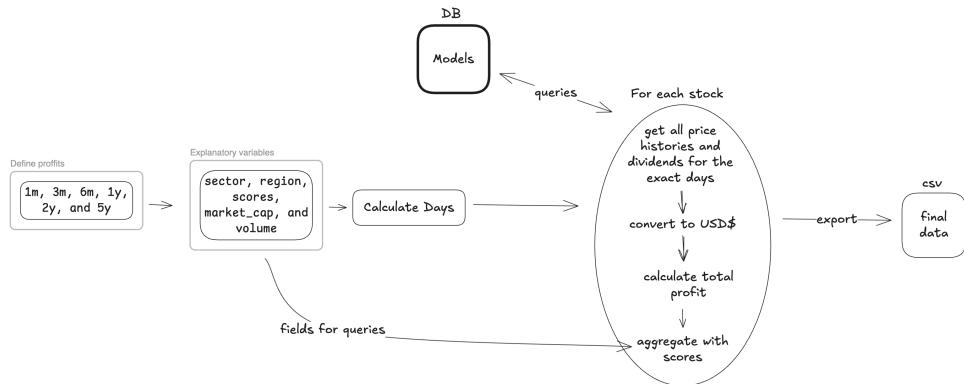


Figure 3.6: First Data Aggregation Schema

Once we had the final csv file we noticed that we had too many empty values for profits, the problem was because some of the days calculated were empty for a specific company, so there was no price history for them. Since the companies belong to multiple countries, the different holidays could be causing another major loss in the information. To fix this in a general way we implemented an internal method for estimating price histories if they didn't exist for the wanted date. The complete implementation of this data export and price estimation process can be found in Appendix 7.3.

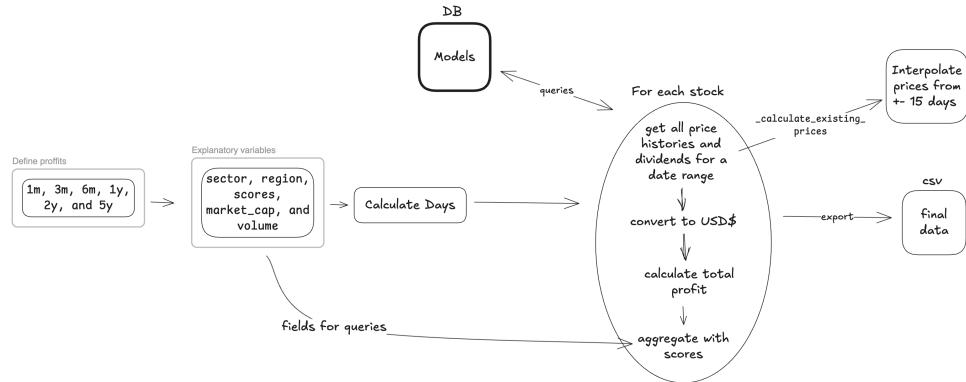


Figure 3.7: Final Data Aggregation Schema

This can be done due to the fact that we are only looking for long term profitabilities, so we can **approximate the price** on a day  $x$  by interpolating the price on  $x - y$  and  $x + z$ , weighted by how close they are to the original day  $x$ , with a maximum distance from  $x$  of 15 days. This significantly reduced the amount of missing values.

### 3.2.2 Datasets Construction

For this study, we have attacked the main objective from multiple perspectives. First we made a descriptive analysis to understand the data, and then we proposed different models to see if there were any relationships between the scores factors and long term profitabilities. Thats why when creating the datasets, we made two different approaches to have more information for the analysis:

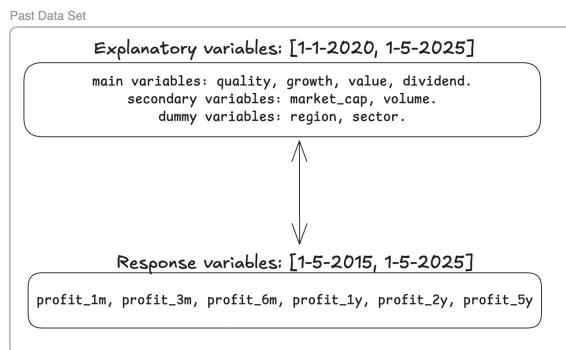


Figure 3.8: Past Dataset

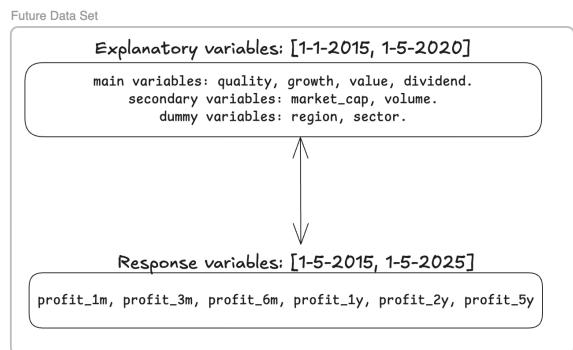


Figure 3.9: Future Dataset

For the past dataset, we calculated the scores for the range 2020-2025 and then related them to the profits obtained until the specified date. For example, the date 1-1-2021 would be linked to the 2 years' profitability obtained until that date:

$$Profit_{1-1-2021}(\%) = \frac{P_{1-1-2021} - P_{1-1-2019} + D_{acc}}{P_{1-1-2019}} \quad (3.1)$$

Which interprets to **how earlier profits have influenced on the explanatory variables**. This will give us valuable information in the descriptive analysis.

For the model fitting, **we want to know if the current scores are good predictors of the future profits**. So we calculated the scores for the range 2015-2020 and related them to the profits obtained from that date to the future. For example, the date 1-1-2021 would be linked to the 2 years' profitability that will be obtained in the future:

$$\text{Profit}_{1-1-2021}(\%) = \frac{P_{1-1-2023} - P_{1-1-2021} + D_{\text{acc}}}{P_{1-1-2021}} \quad (3.2)$$

As this approach is what we actually want to use for the predictive models, we will be using the future dataset structure for most of the analysis.

### 3.2.3 Preprocessing

Even with all of these protocols to create the most reliable and filled datasets, we still had to follow the methodologies explained in the previous chapter to analyze the data.

#### Cleaning the Data

To begin, we looked at the data structure and saw that **about 50% of growth scores were empty**. This turned on many alerts from problems with the scores calculation algorithm, because compared to the rest of the variables there was a significant difference, since the **normal absence was only around 8%** taking into account the value score as reference because it doesn't have 10 year averages.

quality	growth	value	profit_1m	profit_3m	profit_6m	profit_1y	profit_2y	profit_5y
13.4	54.0	7.4	2.6	3.1	3.2	3.2	3.3	6.0

**Table 3.3:** Missing values (%) per column for Small Data Future.

But after digging into the data we saw that it was due to the fact that many companies stopped sending reports or selling but kept existing, so we decided to deleted them from our dataset because they weren't behaving as a "normal company", and this could create a bias in the models. Additionally, as noted earlier, we replaced all NaN values in the dividend scores with 0, since Tweenvest's algorithm omits a score when a company does not pay dividends. That's why we don't see any missing values in the dividend scores.

### 3.2.4 Outlier Detection

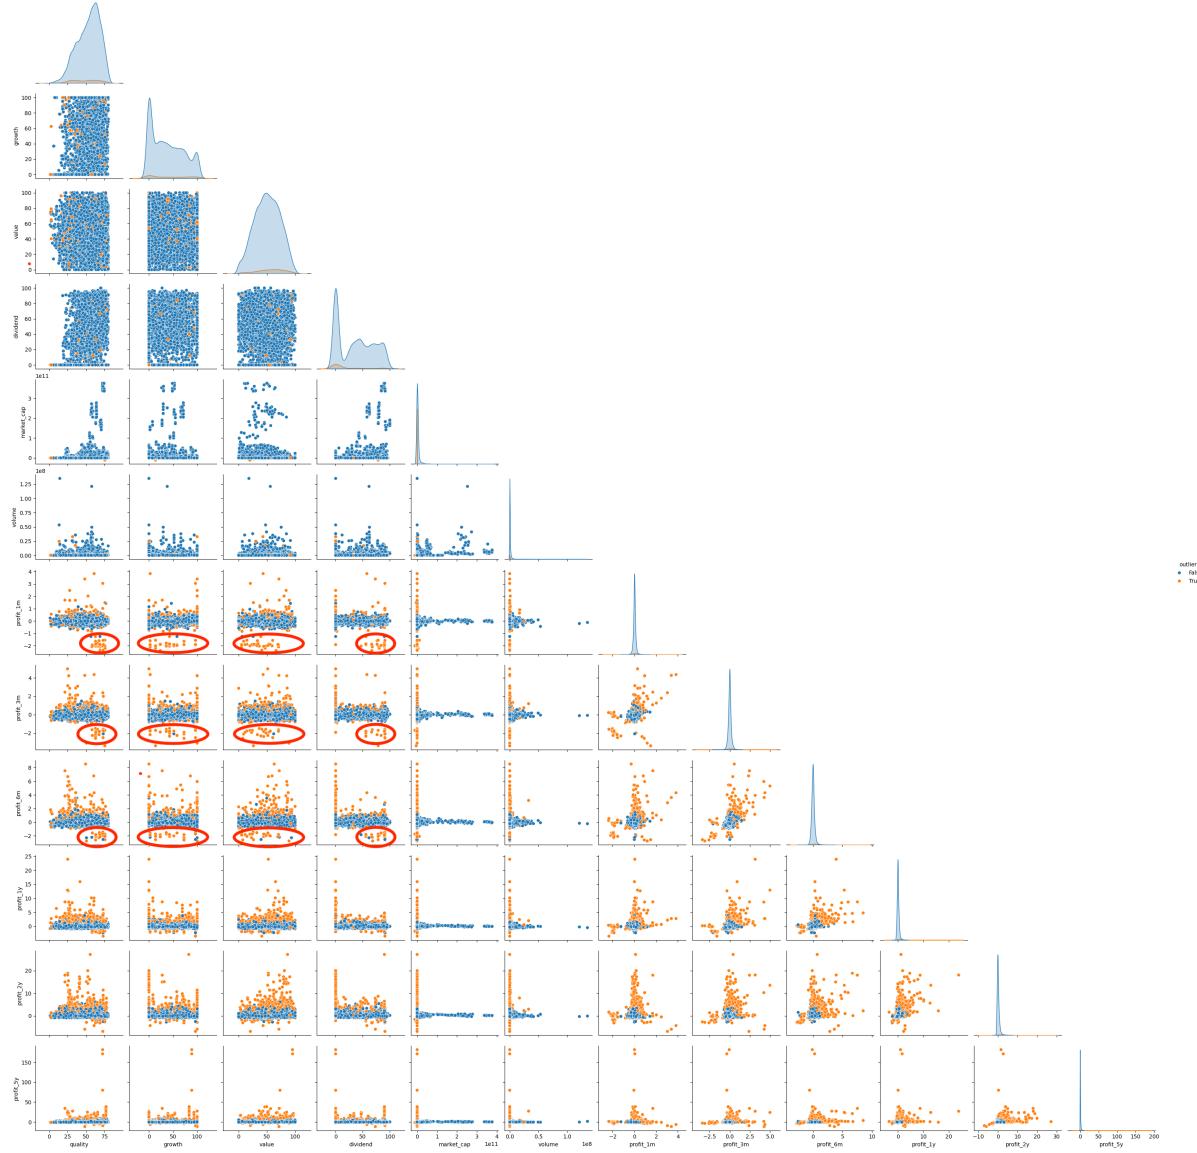
In this section we will present the performance of the different outlier detection methods. But the main analysis of the final datasets will be done in the results chapter.

#### Inter Quartile Range

This method is usefull for big datasets that need to be filtered quickly and where the meaning of the outliers is not multivariate. But as it will be explained in the results chapter, we saw correlations between the different profits, so we needed to use multivariate outlier detection methods, otherwise we could be creating a bias in the data for the model fitting.

## Isolation Forest

We started of using this technique due to its speed and accuracy for non-linear multivariate clusters.

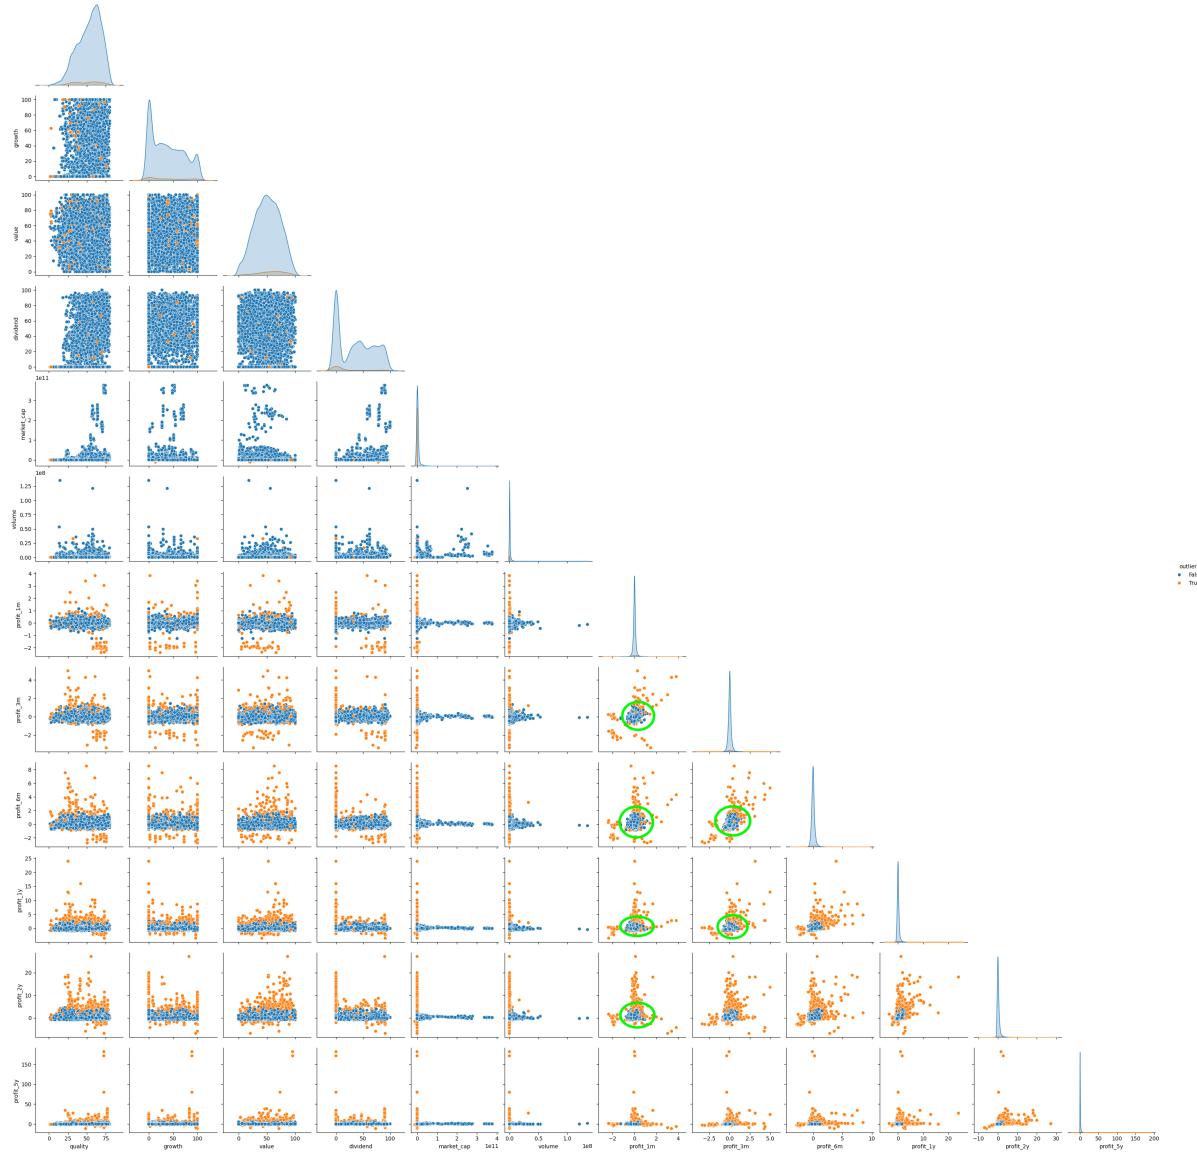


**Figure 3.10:** Small Data Future - IF (5% tolerance)

This pairplot shows how the IF method is detecting outliers in the profits data. In this case it is detecting as outliers much of the data that seems to form a structure in the profits. Also, it is deleting some interesting clusters when comparing the profits with the scores, and here we might lose some important information– this is highlighted with the red circles.

### Single Vector Machine

When using the SVM we saw a similar detection compared to the IF method, which is very interesting since their approach is completely different.

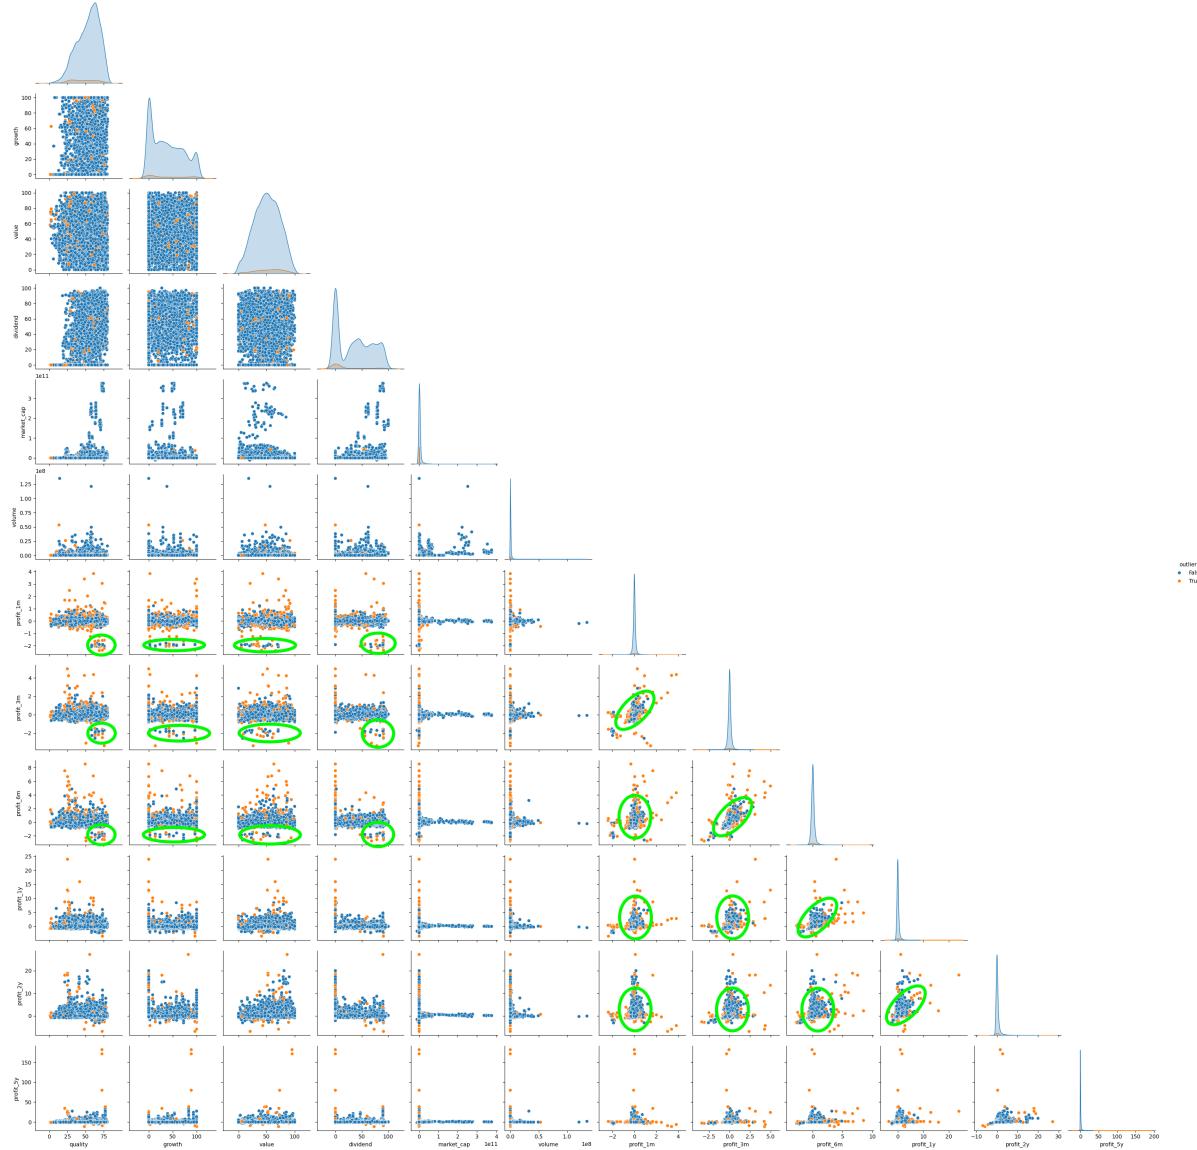


**Figure 3.11:** Small Data Future - SVM (5% tolerance)

But in this case also see a slight improvement with recuperating the structures in some of the profits– this is circled with the green color.

### Local Outlier Factor

Surprisingly, with the LOF method we see a good improvement for both the profits and the scores clusters– which is highlighted with the green circles.

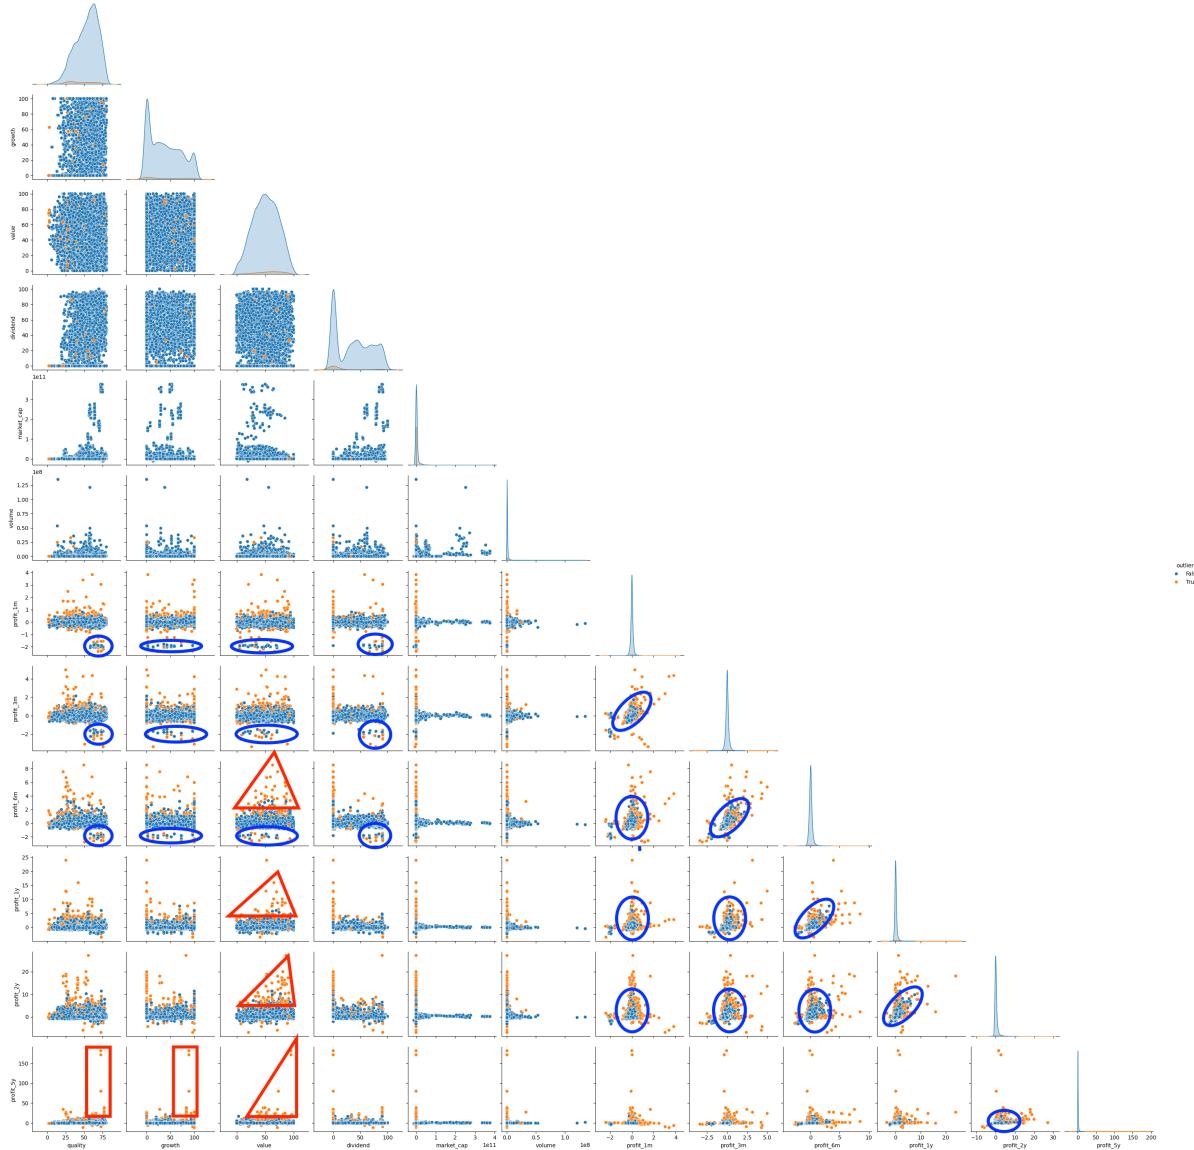


**Figure 3.12:** Small Data Future - LOF (5% tolerance)

We knew we were going in the right direction, but with this method we were missing some outliers that were detected before in the score-profits clusters.

## Multi-Criteria Outlier Detection

That's when we realized that if each method was good at detecting some outliers but bad at others, it makes a lot of sense to use them together. Following this approach and after multiple tests, we found that the best results came when increasing the overall tolerance and then selecting the intersection of outliers between them, so with a tolerance of 20% we only excluded around 4% of the data.



**Figure 3.13:** Small Data Future - Multi-Criteria (20% tolerance)

As shown in the pairplot, this method is able to maintain most of the structures of the profits and also keep the small clusters for the scores, which is highlighted with the blue circles. But of course, it is not perfect.

For future developments of these type of techniques, we have to mention that understanding the main core ideas of each method we could probably optimize the results by changing the tolerance parameters individually by applying multipliers for each method:

1. SVM is basically creating geometrical hyper-dimensional frontiers to form a safe hyper-volume of data clusters. So the intuition tells us that this method should probably have a

**higher tolerance** than the rest.

2. LOF is looking for clusters of high density, not taking account the distance between the clusters. So this method is very good for detecting extreme phenomena but that could be consistent. Its **tolerance could be between the other methods**.
3. IF is looking for the most isolated points following a random forest approach. It is good for detecting those points that seem to make not much sense, so its **tolerance shouldn't be as high**.

But this is just a hypothesis, we will have to be studied it in the future with rigorous tests.

Finally, we have to clarify that outliers are anomalies in the behavior of the data, so even though we are excluding them for obtaining a general set of models. We could be missing valuable information, such as the possible structures shown with the red figures. Thats why we propose a separate analysis for the outliers, to see if really good profits (or bad ones) have any characteristics, such as average high or low scores.

### 3.2.5 Summary

To summarize how we have treated the data, we ended up having 9 different datasets, each with 17 columns where we have the explained variables, and also the company's name and stock ticker (see Appendix Figure 7.1).

Dataset	Raw	Cleaned	No Outliers
<b>Big Data Past<sup>†</sup></b>	16,119,852	9,991,895	9,492,300
<b>Big Data Future<sup>†</sup></b>	31,933,137	14,518,624	13,792,692
<b>Small Data Past</b>	83,243	48,338	46,415
<b>Small Data Future</b>	65,166	26,214	25,220
<b>Small Data Future USA</b>	60,503	21,240	20,463
<b>Small Data Future EU</b>	62,725	25,217	24,245
<b>Small Data Future AS</b>	65,772	28,085	27,156
<b>Small Data Future LAT</b>	48,412	20,610	19,851
<b>Small Data Future AF</b>	58,341	21,182	20,335

**Table 3.4:** Rows per Dataset

<sup>†</sup> In the *Big Data Past* and *Big Data Future* datasets, we removed the outliers using the Isolation Forest method for saving computational time.



# Chapter 4

## Results

### 4.1 Descriptive Analysis

Before jumping into the model fitting, it is necessary to understand the data that we are working with.

#### 4.1.1 Variables Distributions

##### Scores

Following the scores order, quality's distribution show a **clear negative skewness**, meaning that there are more companies closer to having a good quality score. This could be due to the fact that we have deleted companies that don't behave as a "normal company", so we will analize the distributions of the profits to see if there is a bias in the data. Also, there is a **secondary peak around the 30-40 scores**, which could be related to the time-frame of the data.




Figure 4.1: Big Data Future - Scores




Figure 4.2: Big Data Future - Scores Deep

Around 2015-2020 interest rates were at a historical low, so there could be more companies with more debt than normal. Some central banks, particularly in Europe, experimented with negative interest rate policies on deposit facilities during this period (Figure 4.3). This possibility should be studied in more detail, but knowing that this doesn't happen for the dataset between

2020-2025 where interest rates are higher, we have a clue that this could be what's causing it (see Appendix Figure 7.3).

### CHART EURIBOR 1 MONTH

Historical Euribor rates



**Figure 4.3: EMMI - Euribor Interest Rates**

Finishing with quality, there is a "defect" in the dataset. We don't have 10 year averages for the companies that are in the 2015-2020 period – Tweenvest is only paying for data after 2012– so the algorithm gives 0 to some of its criterias making the maximum quality be 80. But since this is happening for all of the companies, **it is not a problem for the model fitting**, although we won't be analyzing those criteria's performance.

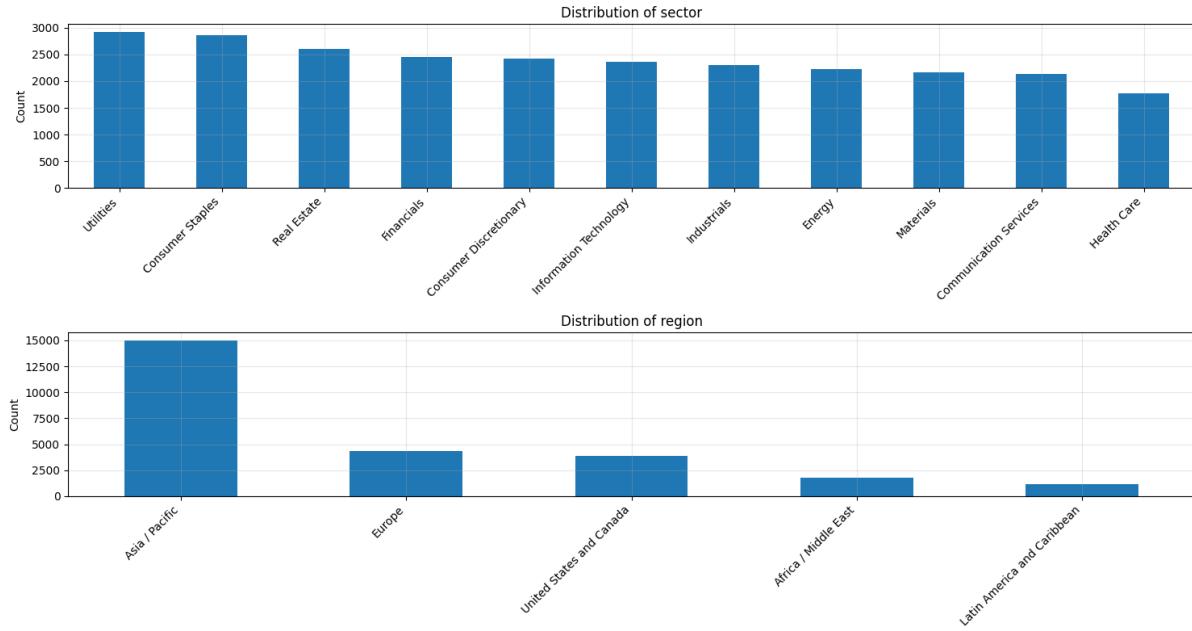
Continuing with growth, we notice that there are many companies with a score of 0 and a score of 100, and something similar happens with the 0s in Dividend. For having a deeper look, we excluded those values for the Figure 4.2 where we can see that it is closer to a uniform distribution with a descend slope after 65, which means that most companies have a similar growth rate, but fewer have higher ones.

When looking at the value score, we can see that the distribution almost has an average of 50, and there is a continuum of values through the whole range, so it seems that we have a well scale for distributing the companies.

Lastly for the dividend score, as mentioned before, there are many companies that don't pay dividends. But aside from that, we have an irregular distribution for the score, with a peak around 85-92 meaning that there are multiple companies that respect their dividend policy.

### Dummies

When it comes to the dummies we see a uniform distribution with the sectors because this condition was forced in the data creation process. But with the regions we have a different result, seeing a clear **predominance of the Asiatic region**, which happens in all of the datasets. Which could create issues for the model fitting, since it is known that in countries like China have more volatility and lower returns Chen et al. [3].

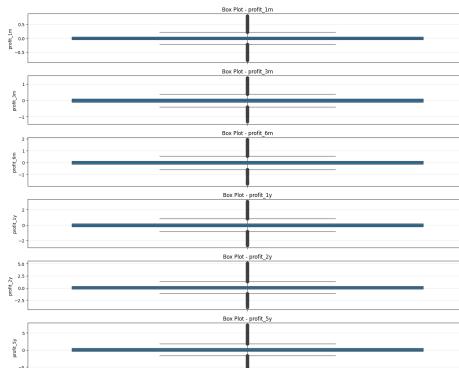

**Figure 4.4:** Small Data Future - Dummies

Also, their auditors often lack true independence, especially when auditing state-owned enterprises (SOEs) or firms with strong political ties. The Chinese government tightly controls domestic accounting firms and imposes restrictions on foreign auditors. Even international firms like *PwC-China* operate under local joint ventures, often with limited autonomy, as reviewed in Liu and Tian [11]. So we will have to check if the dummies variables are enough to take this effect into account, or if we need a different approach for modeling the geographical effects.

It is worth noting that the distribution of companies across sectors varies significantly over time, as shown in Appendix Figures 7.6 and 7.7. This temporal variation in sector representation presents an interesting opportunity for future research.

## Profits

When analyzing the profit distributions across different time horizons, we observe that the distributions contain a numerous amount of outliers that don't let us see the distributions clearly (Appendix Figures 7.8 and 7.9). So just for this descriptive analysis purposes, **we used the IQR method** mentioned in previous chapters to remove the outliers.

**Figure 4.5:** Big Data Future IQR - Profits Boxplot**Figure 4.6:** Big Data Future IQR - Profits

Interestingly, we notice that the short-term profit distributions (1-month and 3-month) are more symmetric and closer to normal distributions, which is consistent with the efficient market hypothesis suggesting that short-term price movements are largely random.

However, as we extend to longer time horizons, the distributions become more positively skewed, which aligns with the general expectation that most companies don't have good returns, except of the market winners that tend to generate extraordinary positive returns due to the mentioned **economic moats**. This is particularly evident in the longer-term profit distributions (2-year and 5-year), where we see a longer tail extending toward higher positive returns.

There is also a tendency to having less negative returns compared to the positive ones the longer the time-frame. Logically, companies with bad returns don't last long.

### 4.1.2 Correlations

For analyzing the correlations between the variables, we will start firstly with just the cleaned big datasets, and then we will use also the datasets without the outliers, since these are the ones that will be used for the model fitting.

Here is the only section that we will be using the past Datasets, to see if past profits are being reflected in current scores.

#### Clean Datasets

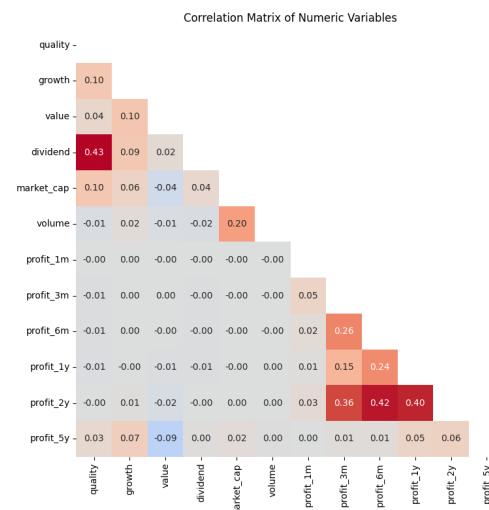


Figure 4.7: Big Data Past - Correlations





Figure 4.8: Big Data Future - Correlations

We can appreciate three main expected correlations:


- **Companies with good quality** –that is so, they have consistent growth, good liquidity and healthy debt– are more likely to give dividends.
- There is a strong relationship between the market cap of a company and the volume of the shares traded.
- **Companies have inertia**, or how it is called *Momentum*, meaning that the profit of a company is positively correlated with the profit of the previous period.

## Outliers Datasets

But as it has been seen multiple times through the analysis, the outliers create a big bias that doesn't allow us to see the real correlations occurring for "normal companies". So as shown in the table 3.4, we will use the Big datasets filtered with the Isolation Forest method to remove outliers as a first approach, but for the model fitting we will use the multi-criteria outlier detection.



**Figure 4.9:** Big Data Past IF - Correlations



**Figure 4.10:** Big Data Future IF - Correlations

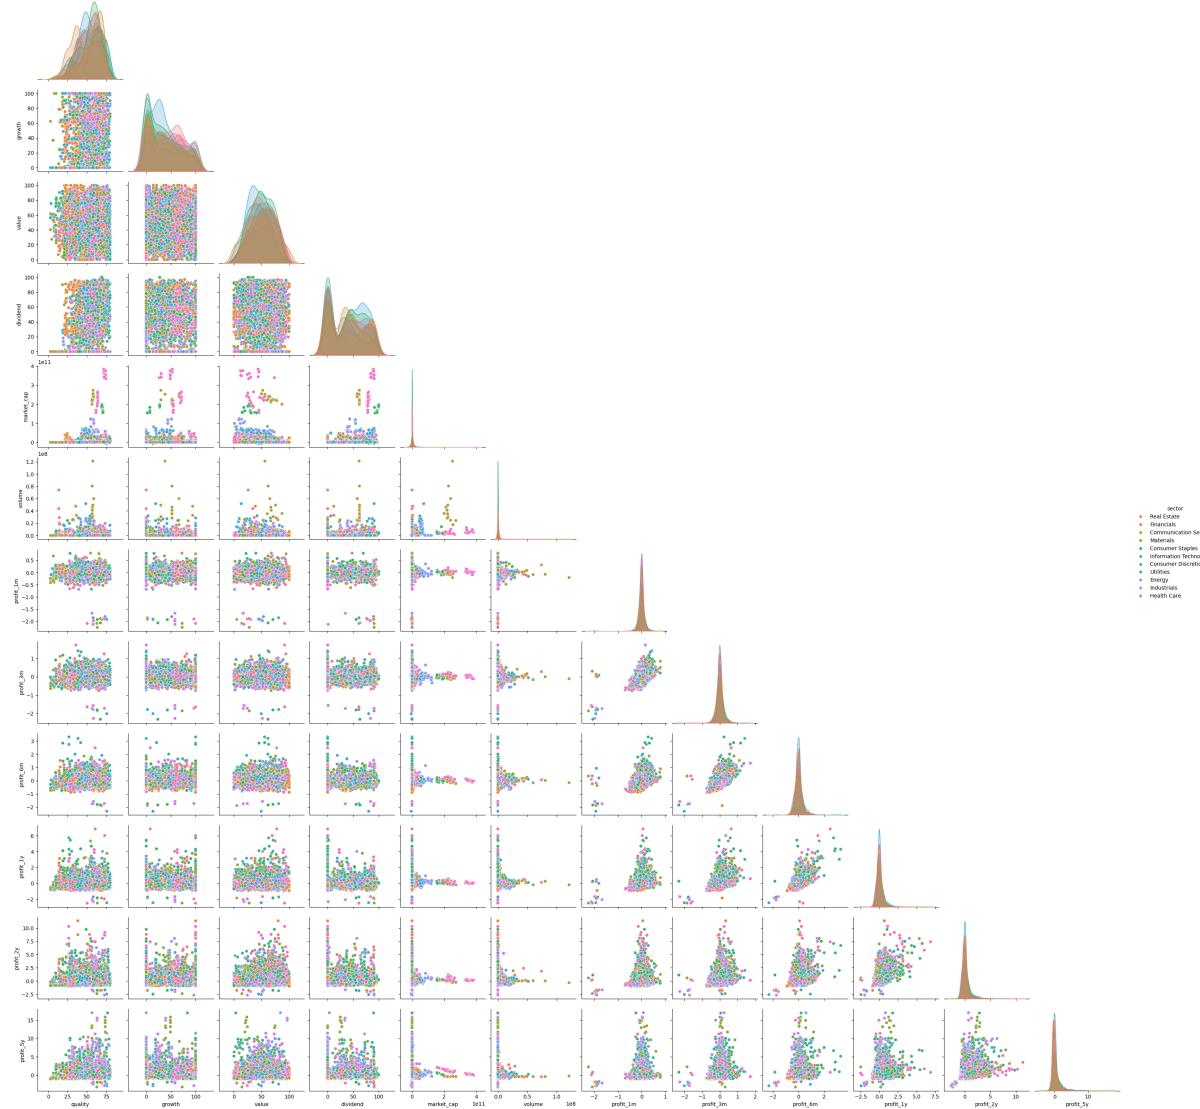
As expected, the correlations already seen stay similar with a slight increase in the *Momentum*, showing more importance of it than other variables. But actually, there are **new correlations seen** for both datasets:

- **Accumulated profits affect the scores.** In the past dataset we can see that long term past profits translate into higher quality and dividend scores– as expected by their correlation. But also, they make the value score to decrease. This could be explained by two effects:
  - **Bull markets.** As explained in the second chapter, a company can experience good returns in the stock market that are not followed by good financial growth. This makes their value multiples to raise, making their value score to decrease.
  - **The company is not able to maintain its profitability.** If a company has had good past profits, but is not able to maintain them, the market will eventually realize that and the company will have a lower value score.
- **Long term stock returns affect the company's growth.** This is a very interesting result, since it explains the reason why companies try to go to the market to raise capital and help the company grow.
- **Quality is correlated to future profits.** This could be explained by the fact that a company with good quality is more likely to have good future profits, but also by the fact that the market is efficient and the price of a company will reflect all of the information available.
- There is a **mild correlation between most scores**, which could be explain by the fact that a company that does well in one aspect, is more likely to do well in the others. This correlation won't affect the models because they are separated by score.

Althought we have to make clear that these correlations do not stand along all of the datasets separated by region (see Appendix Figures from 7.10 to 7.19), so we aren't sure yet if these relationships will be properly explained with linear models.

### 4.1.3 Data clusters

Finally, we are going to try to compare these results with the pairplots of the datasets, to look for clusters of data when comparing many variables at once.



**Figure 4.11:** Pairplot of Small Data Future - Multi

Here it is harder to eye-see the data structure that could form clusters since there are many datapoints, but there are some exceptions:

- **Profits.** We can see linear correlations between some of the different profits, but for longer horizons the relationships are no longer liner. Which is expected since the further the horizon, more variables take part in the calculation of the profit.
- **Market Cap.** Companies with bigger market caps tend to have higher quality and dividend scores. Bigger companies should have a more stable financial situation, otherwise they wouldn't be as big in the long term. Also there are some sectors that tend to have higher market caps than others, like in Health Care or Consumer Staples.
- **Sector.** It seems that there could be a pattern seen in the sectors distribution for each scores, but it is not clear.

Althought, these results are not as clear and they also variate deppending on the dataset. Showing early results of non-linear relationships with the dummy variables.

## 4.2 Predictive Models

### 4.2.1 Linear Regression

### 4.2.2 The problem with ARIMA

### 4.2.3 GAM

### 4.2.4 Neural Network

## 4.3 Back-testing



## **Chapter 5**

# **Discussion**



## **Chapter 6**

## **Conclusion**



# Bibliography

- [1] A. A. Abro, E. Taşci, and A. Uğur. "A Stacking-based Ensemble Learning Method for Outlier Detection". In: *Balkan Journal of Electrical & Computer Engineering* 8.2 (Apr. 2020). URL: <https://dergipark.org.tr/en/download/article-file/1106223>.
- [2] Markus M. Breunig et al. "LOF: identifying density-based local outliers". In: *ACM sigmod record*. 2000. URL: <https://dl.acm.org/doi/pdf/10.1145/335191.335388>.
- [3] X. Chen et al. "Economic policy uncertainty and the predictability of stock returns: impact of China in Asia". In: *China Accounting and Finance Review* 26.5 (2024), pp. 645–679. DOI: 10.1108/CAFR-11-2023-0144. URL: <https://doi.org/10.1108/CAFR-11-2023-0144>.
- [4] scikit-learn developers. *OneClassSVM*. scikit-learn Documentation. URL: <https://scikit-learn.org/stable/modules/generated/sklearn.svm.OneClassSVM.html>. 2025.
- [5] Official Documentation. *bulk\_create Method*. Django Documentation. URL: <https://docs.djangoproject.com/en/5.2/ref/models/querysets/#bulk-create>. 2025.
- [6] Official Documentation. *select\_related Method*. Django Documentation. URL: [https://docs.djangoproject.com/en/5.2/ref/models/querysets/#django.db.models.query.QuerySet.select\\_related](https://docs.djangoproject.com/en/5.2/ref/models/querysets/#django.db.models.query.QuerySet.select_related). 2025.
- [7] P. Dorsey and J. Mansueto. *The Five Rules for Successful Stock Investing*. 2011.
- [8] Walter Enders. *Applied Econometric Time Series*. URL: <http://sggscclibrary.saraswatilib.com/UploadImg/637617646888135985.pdf>. 1948, p. 293.
- [9] John D. Hunter et al. *Matplotlib: Python plotting package*. Online documentation. URL: <https://matplotlib.org/stable/>. 2025.
- [10] Fei Tony Liu, Kai Ming Ting, and Zhi-Hua Zhou. "Isolation-based anomaly detection". In: *ACM Transactions on Knowledge Discovery from Data (TKDD)* 6.1 (2012), p. 3. URL: <https://www.lamda.nju.edu.cn/publication/tkdd11.pdf>.
- [11] Qigui Liu and Gary Tian. "Controlling shareholder, expropriations and firm's leverage decision: Evidence from Chinese Non-tradable share reform". In: *Journal of Corporate Finance* 18.4 (2012). Special Section: Contemporary corporate finance research on South America, pp. 782–803. ISSN: 0929-1199. URL: <https://www.sciencedirect.com/science/article/pii/S0929119912000533>.
- [12] MSCI. "Fundamental Data Methodology". In: (June 2025). URL: <https://www-cdn.msci.com/documents/10199/0b41b3a3-9da8-46f6-fce7-1dc1177e4a23>.
- [13] Daniel Servén and Charlie Brummitt. *Generalized Additive Models*. pyGAM Documentation. URL: [https://pygam.readthedocs.io/en/latest/notebooks/tour\\_of\\_pygam.html](https://pygam.readthedocs.io/en/latest/notebooks/tour_of_pygam.html). 2018.
- [14] William F. Sharpe. *Capital Asset Prices: A Theory of Market Equilibrium under Conditions of Risk*. Vol. 19. 3. Sept. 1964, pp. 425–442.
- [15] The pandas development team. *pandas Documentation*. pandas Documentation. URL: <https://pandas.pydata.org/>. 2024.

- [16] Michael L. Waskom. *seaborn: statistical data visualization*. Seaborn Documentation. URL: <https://seaborn.pydata.org/index.html>. 2024.
- [17] Noumir Zineb, Paul Honeine, and Cedue Richard. “On simple one-class classification methods”. In: *IEEE International Symposium on Information Theory Proceedings* (2012). URL: [https://www.researchgate.net/publication/261248728\\_On\\_simple\\_one-class\\_classification\\_methods](https://www.researchgate.net/publication/261248728_On_simple_one-class_classification_methods).

# **Part II**

# **Appendix**



# Chapter 7

## Code Listings

### 7.1 Historical Scores Job Implementation

The following code snippets shows the complete implementation of the historical scores job that was used to populate the database with historical factor scores data:

#### 7.1.1 Enqueueing Historical Scores Job

```
1 @job("historical_scores", timeout="3h")
2 def enqueue_calculate_index_scores_task(
3     start_date: datetime.date,
4     end_date: datetime.date,
5     index_names: list[str] | None = None,
6     region_names: list[str] | None = None,
7     total_stocks: int | None = None,
8     seed: float | None = None,
9     daily: bool = False,
10):
11     logger.info(
12         "[H1] Starting enqueue_year_tasks from %s to %s",
13         start_date,
14         end_date,
15     )
16
17     days = get_range_from_dates(
18         start_date,
19         end_date,
20         delta=relativedelta(days=1) if daily else relativedelta(
21             months=1),
22     )
23     days.reverse()
24
25     if not index_names:
26         index_names = list(
27             Index.objects.exclude(name__in=[Index.INDEX_ALL, "Other"])
28             .only("name")
29             .values_list("name", flat=True)
30         )
31     logger.info("[H1] Enqueuing tasks for %s indexes", len(index_
32         names))
33
34     stocks_per_index = (
35         total_stocks // len(index_names)
36         if total_stocks
37         else 1
38     )
39
40     for day in days:
41         for index_name in index_names:
42             stocks_to_enqueue = stocks_per_index
43             if total_stocks < stocks_to_enqueue * len(index_names):
44                 stocks_to_enqueue = total_stocks
45
46             enqueue_index_scores_task(
47                 day,
48                 index_name,
49                 stocks_to_enqueue,
50                 seed=seed,
51                 daily=daily,
52             )
53
54     logger.info("[H1] All tasks enqueued, waiting for them to finish")
55
56     return True
```

```

33         int(total_stocks / len(index_names)) if total_stocks else
34             None
35
36     stock_count = 0
37     total_stocks_dict = {}
38     for index_name in index_names:
39         stock_ids = get_stock_ids_for_index(
40             index_name,
41             stocks_per_index,
42             start_date,
43             end_date,
44             seed,
45             region_names,
46         )
47         if not stock_ids:
48             raise ValueError(
49                 f"[H1] No stocks found for index '{index_name}' after
50                     {start_date}"
51             )
52         total_stocks_dict[index_name] = stock_ids
53     for day in days:
54         calculate_index_scores_task.delay(
55             index_name=index_name,
56             day=day,
57             stock_ids=stock_ids,
58         )
59         stock_count += len(stock_ids)
60
61     logger.info(
62         "[H1] Enqueued days %s for stocks:%s .", len(days), stock_
63             count
64     )
65
66     return {
67         "stocks_processed": stock_count,
68         "days": len(days),
69         "total_stocks": total_stocks_dict,
70     }

```

**Listing 7.1:** Enqueueing Historical Scores Job

This job function is responsible for:

- Accepting date ranges and configuration parameters for historical score calculation
- Generating a list of dates to process (either daily or monthly intervals)
- Retrieving stock IDs for each index based on the specified criteria
- Enqueuing individual calculation tasks for each day and index combination
- Providing detailed logging and return statistics about the processing

## 7.1.2 Calculate Index Scores Data

```

1
2 @job("historical_scores", timeout="45m")
3 def calculate_index_scores_task(
4     index_name: str,
5     day: datetime.date,
6     stock_ids: list[int],
7 ):
8     start_time = time.perf_counter()
9     logger.info(
10         "[H2] Calculating index",
11         extra={
12             "index": index_name,
13             "day": day,
14             "stocks": len(stock_ids),
15         },
16     )
17
18     # ----- Index calculation ----- #
19     index_stats = get_indexes_information(
20         index_names=[index_name], calculation_date=day
21     )
22     index_quality = get_quality_ratios_by_index(
23         index_names=[index_name], calculation_date=day
24     )
25     stats = index_stats.get(index_name)
26     quality = index_quality.get(index_name)
27
28     # -----
29     if (
30         not stats
31         or not quality
32         or all(value is None for value in quality.values())
33         or all(
34             value is np.nan
35             for stat in stats.values()
36             for value in stat.values()
37         )
38     ):
39         raise RuntimeError(
40             f"[H2] Missing index data for {index_name} on {day},
41             stats: {stats}, quality: {quality}"
42         )
43
44     calculate_stocks_scores_task.delay(
45         day=day,
46         stats=stats,
47         quality=quality,
48         stock_ids=stock_ids,
49     )
50     return {
51         "index_calculated": index_name,
52         "total_time": time.perf_counter() - start_time,
53     }

```

Listing 7.2: Calculate Index Scores Data

This job function is responsible for:

- Calculating all the necessary index stats data
- Enqueuing the calculation of the stocks scores for the given index and date
- Providing detailed logging and return statistics about the processing

### 7.1.3 Calculate Stocks Scores Task

```

1 @tracer.start_as_current_span("calculate_stocks_scores_task")
2 @job("historical_scores", timeout="4h")
3 def calculate_stocks_scores_task(
4     day: datetime.date,
5     stats: dict,
6     quality: dict,
7     stock_ids: list[int],
8 ):
9     """
10     Job 3: For each stock, calculate scores and bulk save.
11     """
12     start_time = time.perf_counter()
13     logger.info(
14         "[H3] Calculating scores",
15         extra={"day": day, "stocks": len(stock_ids)},
16     )
17
18     if not stock_ids:
19         return {
20             "stocks_processed": 0,
21         }
22     # Select related for trying to avoid multiple queries to the DB
23     stocks = Stock.objects.filter(id__in=stock_ids).select_related(
24         "share_type",
25         "share_type__company",
26         "share_type__company__industry",
27         "share_type__company__industry__industry_group",
28         "share_type__company__industry__industry_group__sector",
29     )
30     factor_scores = []
31     for stock in stocks:
32         score = calculate_scores(
33             stock,
34             stats,
35             quality,
36             calculation_date=day,
37             bulk_create_or_update=True,
38         )
39         if not score:
40             logger.warning(
41                 "[H3] No score for calculated",
42                 extra={"stock_id": stock.pk, "day": day, "score": score},
43             )
44             continue # Not as important as the others, need to check
45             errors in DataDog
46         factor_scores.append(score)
47
48     bulk_create_or_update(
49         model=FactorScore,

```

```

49     object_list=factor_scores,
50     unique_fields=["stock", "date"],
51     update_fields=["quality", "growth", "value", "dividend"],
52 )
53
54     return {
55         "stocks_processed": len(stock_ids),
56         "total_time": time.perf_counter() - start_time,
57         "time_per_stock": (time.perf_counter() - start_time) / len(
58             stock_ids),
59     }

```

**Listing 7.3:** Calculate Stocks Scores Task

This job function is responsible for:

- Calculating the factor scores for each stock in the given index and date
- Bulk saving the factor scores for the given index and date
- Providing detailed logging and return statistics about the processing

## 7.2 Workers Management

In this section we have 2 different configurations for the workers management, since we had to switch servers.

### 7.2.1 Tweenvest's Production Server

```

1 [program:worker1]
2 command=python manage.py rqworker internal_tasks ciq_priority fmp_
    priority ciq indices_partial indices_full historical_scores fmp_
    emails --with-scheduler
3 directory=/app
4 autostart=true
5 autorestart=true
6 redirect_stderr=true
7 stdout_logfile=/dev/fd/1
8 stdout_logfile_maxbytes=0
9 environment = SERVICE_NAME="tweenvest-worker"
10
11 [program:worker2]
12 command=python manage.py rqworker internal_tasks ciq_priority fmp_
    priority ciq indices_partial indices_full historical_scores fmp_
    emails --with-scheduler
13 autostart=true
14 directory=/app
15 autorestart=true
16 redirect_stderr=true
17 stdout_logfile=/dev/fd/1
18 stdout_logfile_maxbytes=0
19 environment = SERVICE_NAME="tweenvest-worker"
20
21 [program:worker3]
22 command=python manage.py rqworker internal_tasks ciq_priority fmp_
    priority ciq indices_partial fmp emails historical_scores --with-
    scheduler

```

```
23 autostart=true
24 directory=/app
25 autorestart=true
26 redirect_stderr=true
27 stdout_logfile=/dev/fd/1
28 stdout_logfile_maxbytes=0
29 environment = SERVICE_NAME="tweenvest-worker"
30
31 [program:worker4]
32 command=python manage.py rqworker ciq_priority fmp_priority ciq
            indices_partial fmp emails historical_scores --with-scheduler
33 autostart=true
34 directory=/app
35 autorestart=true
36 redirect_stderr=true
37 stdout_logfile=/dev/fd/1
38 stdout_logfile_maxbytes=0
39 environment = SERVICE_NAME="tweenvest-worker"
40
41 [program:worker5]
42 command=python manage.py rqworker fmp_priority ciq_priority ciq fmp
            emails --with-scheduler
43 autostart=true
44 directory=/app
45 autorestart=true
46 redirect_stderr=true
47 stdout_logfile=/dev/fd/1
48 stdout_logfile_maxbytes=0
49 environment = SERVICE_NAME="tweenvest-worker"
50
51 [program:worker6]
52 command=python manage.py rqworker emails --with-scheduler
53 autostart=true
54 directory=/app
55 autorestart=true
56 redirect_stderr=true
57 stdout_logfile=/dev/fd/1
58 stdout_logfile_maxbytes=0
59 environment = SERVICE_NAME="tweenvest-worker"
60
61 [program:worker7]
62 command=python manage.py rqworker ciq_priority ciq emails summaries
            --with-scheduler
63 directory=/app
64 autostart=true
65 autorestart=true
66 redirect_stderr=true
67 stdout_logfile=/dev/fd/1
68 stdout_logfile_maxbytes=0
69 environment = SERVICE_NAME="tweenvest-worker"
70
71 [program:worker8]
72 command=python manage.py rqworker ciq_priority ciq emails summaries
            --with-scheduler
73 directory=/app
74 autostart=true
75 autorestart=true
76 redirect_stderr=true
```

```

77 stdout_logfile=/dev/fd/1
78 stdout_logfile_maxbytes=0
79 environment = SERVICE_NAME="tweenvest-worker"
80
81 [program:worker9]
82 command=python manage.py rqworker historical_scores ciq_priority
     summaries ciq emails --with-scheduler
83 directory=/app
84 autostart=true
85 autorestart=true
86 redirect_stderr=true
87 stdout_logfile=/dev/fd/1
88 stdout_logfile_maxbytes=0
89 environment = SERVICE_NAME="tweenvest-worker"
90
91 [program:worker10]
92 command=python manage.py rqworker summaries ciq_priority fmp_priority
     ciq fmp emails --with-scheduler
93 directory=/app
94 autostart=true
95 autorestart=true
96 redirect_stderr=true
97 stdout_logfile=/dev/fd/1
98 stdout_logfile_maxbytes=0
99 environment = SERVICE_NAME="tweenvest-worker"

```

Listing 7.4: Production Workers Management

## 7.2.2 Personal Development Server

```

1 [program:worker1]
2 command=python manage.py rqworker historical_scores
3 process_name=%(program_name)s_%(process_num)02d
4 numprocs=16
5 directory=/app
6 autostart=true
7 autorestart=true
8 redirect_stderr=true
9 stdout_logfile=/dev/fd/1
10 stdout_logfile_maxbytes=0
11 environment = SERVICE_NAME="tweenvest-worker"

```

Listing 7.5: Hetzner Workers Management

## 7.3 Data Export Job Implementation

The following code shows the complete implementation of the data export job that was used to create the final dataset, including the price estimation for missing values due to weekends and holidays:

```

1 def export_factors_and_pricehistory_task(
2     stocks_id: list[int],
3     start_date: datetime.date,
4     end_date: datetime.date,
5     export_name: str,

```

```
6     daily: bool = False,
7 ):
8     """
9     Export FactorScore, PriceHistory y rentabilidades a CSV sin
10    cargar todo en memoria.
11 """
12    start_time = time.perf_counter()
13    print("Starting CSV export...")
14
15    periods = {
16        "profit_1m": relativedelta(months=1),
17        "profit_3m": relativedelta(months=3),
18        "profit_6m": relativedelta(months=6),
19        "profit_1y": relativedelta(years=1),
20        "profit_2y": relativedelta(years=2),
21        "profit_5y": relativedelta(years=5),
22    }
23
24    factor_fields = [
25        "stock__ticker",
26        "stock__share_type__company__name",
27        "stock__share_type__company__industry__industry_group__sector
28            __name",
29        "stock__share_type__company__country__region",
30        "date",
31        "quality",
32        "growth",
33        "value",
34        "dividend",
35    ]
36    price_fields = ["market_cap_usd", "volume"]
37
38    days = get_range_from_dates(
39        start_date,
40        end_date,
41        delta=relativedelta(days=1) if daily else relativedelta(
42            months=1),
43    )
44    days.reverse()
45
46    profitability_fields = list(periods.keys())
47    export_fields = factor_fields + price_fields + profitability_
48        fields
49    missing_prices = 0
50    missing_factors = 0
51    fs_path = f"api/factors/data_exports/{export_name}.csv"
52    with open(fs_path, mode="w", newline="", encoding="utf-8") as f:
53        writer = csv.DictWriter(f, fieldnames=export_fields)
54        writer.writeheader()
55
56        for stock_id in tqdm(stocks_id, desc="Stocks"):
57
58            price_qs = PriceHistory.objects.filter(
59                stock__id=stock_id,
60                date__range=(
61                    start_date,
62                    end_date + relativedelta(years=5),
63                ),
64            ).
```

```

60     ).values("date", "close_price", "market_cap", "volume", "fx_mult")
61     price_lookup_stock = {
62         (stock_id, ph["date"]): ph
63         for ph in price_qs
64         if ph["fx_mult"] is not None
65     }
66
67     dividend_qs_stock = DividendHistory.objects.filter(
68         stock_id=stock_id,
69         ex_dividend_date__range=(start_date, end_date),
70     ).values("ex_dividend_date", "adjusted_dividend", "fx_mult")
71     dividend_lookup = {
72         (stock_id, dh["ex_dividend_date"]): (
73             float(dh["adjusted_dividend"]) * float(dh["fx_mult"])
74         )
75         for dh in dividend_qs_stock
76         if dh["fx_mult"] is not None
77     }
78     fs_qs_stock = (
79         FactorScore.objects.filter(stock_id=stock_id, date__in=days)
80         .select_related(
81             "stock__share_type__company__industry__industry_group__sector",
82             "stock__share_type__company__country",
83         )
84         .values(*factor_fields)
85     )
86
87     fs_lookup = { (stock_id, fs["date"]): fs for fs in fs_qs_stock}
88
89     for day in days:
90         factor = fs_lookup.get((stock_id, day))
91         if not factor:
92             missing_factors += 1
93             continue
94         day = factor["date"]
95         price_day = day
96         while price_day.weekday() >= 5:
97             price_day -= datetime.timedelta(days=1)
98
99         price = price_lookup_stock.get((stock_id, price_day))
100        if not price or price["close_price"] is None:
101            price = _calculate_existing_prices(
102                price_day, price_lookup_stock, stock_id
103            )
104            if not price:
105                missing_prices += 1
106                continue
107
108        try:
109            market_cap_usd = (
110                float(price.get("market_cap", 0) or 0)
111                * price["fx_mult"]

```

```

112
113     )
114 except Exception:
115     market_cap_usd = None
116 try:
117     volume = (
118         float(price.get("volume", 0) or 0) * price["
119             fx_mult"]
120     )
121 except Exception:
122     volume = None
123
124 if price["close_price"] is None or price["fx_mult"]
125     is None:
126     continue
127 close_now_usd = float(price["close_price"]) * price["
128     fx_mult"]
129
130 profitabilities = {}
131 for field, delta in periods.items():
132     future_date = day + delta
133     while future_date.weekday() >= 5:
134         future_date -= datetime.timedelta(days=1)
135
136     future_price = price_lookup_stock.get(
137         (stock_id, future_date)
138     )
139     if not future_price or future_price["close_price"]
140         is None:
141         future_price = _calculate_existing_prices(
142             future_date, price_lookup_stock, stock_id
143         )
144     if not future_price:
145         profitabilities[field] = None
146         missing_prices += 1
147         continue
148
149     close_future_usd = float(
150         future_price["close_price"]
151     ) * float(future_price["fx_mult"])
152     if close_future_usd == 0:
153         profitabilities[field] = None
154         missing_prices += 1
155         continue
156     # Dividend sum
157     dividend_sum = sum(
158         v
159         for (s_id, ex_div_date), v in dividend_lookup
160             .items()
161         if s_id == stock_id
162         and day <= ex_div_date < future_date
163     )
164     profitabilities[field] = (
165         close_future_usd + dividend_sum - close_now_
166             usd
167     ) / close_now_usd
168
169     row = {**factor}
170     row["market_cap_usd"] = market_cap_usd

```

```

164         row["volume"] = volume
165         row.update(profitabilities)
166         writer.writerow(row)
167
168     elapsed = time.perf_counter() - start_time
169     print(f"Exportado a {fs_path} en {elapsed:.2f}s")
170     print("MISSING PRICES:", missing_prices)
171     print("MISSING FACTOR SCORES:", missing_factors)

```

**Listing 7.6:** Data Export Job Implementation

This job function is responsible for:

- Exporting factor scores, price history and profitability data to CSV without loading everything into memory
- Handling missing price data due to weekends and holidays by estimating prices from nearby dates
- Converting all monetary values to USD using appropriate exchange rates
- Calculating profitability for different time periods (1m, 3m, 6m, 1y, 2y, 5y)
- Including dividend payments in profitability calculations
- Tracking and reporting missing data points for both prices and factor scores

## 7.4 Custom Python Functions for Data Analysis

```

1 def correlation_matrix(
2     df: pd.DataFrame, columns: list[str], file_path: str | None =
3         None
4 ) -> None:
5     """
6         Plot a correlation matrix for the DataFrame showing only the
7             lower triangle.
8     """
9     df_columns = df[columns]
10    corr_matrix = df_columns.corr()
11
12    # Create a mask for the lower triangle
13    mask = np.triu(np.ones_like(corr_matrix, dtype=bool))
14
15    # Plot using seaborn for better visualization
16    plt.figure(figsize=(12, 8))
17    sns.heatmap(
18        corr_matrix,
19        mask=mask,
20        annot=True,
21        cmap="coolwarm",
22        center=0,
23        fmt=".2f",
24        square=True,
25    )
26    plt.title("Correlation Matrix of Numeric Variables")
27    plt.tight_layout()

```

```

27     if file_path:
28         plt.savefig(file_path)
29         plt.show()
30     else:
31         plt.show()

```

Listing 7.7: Correlation Matrix Plot Function

```

1 def pairplot(
2     df: pd.DataFrame,
3     file_path: str | None = None,
4     kde: bool = False,
5     hue: str = "region",
6     hue_order: list[str] | None = None,
7     palette: dict[str, str] | None = None,
8 ) -> None:
9
10    if len(df) > 10000:
11        df = df.sample(10000, random_state=42)
12
13    g = sns.pairplot(
14        df,
15        hue=hue,
16        diag_kind="kde",
17        hue_order=hue_order,
18        palette=palette,
19        corner=True,
20    )
21    if kde:
22        g.map_lower(sns.kdeplot, levels=4, color=".2")
23
24    if file_path:
25        plt.savefig(file_path)
26        plt.close()
27    else:
28        plt.show()

```

Listing 7.8: Custom Pairplot Function

## 7.5 Data Preprocessing

### 7.5.1 Data Cleaning

```

1 import sys
2 from pathlib import Path
3
4 # Add the parent directory (code/) to sys.path
5 sys.path.append(str(Path().resolve().parent))
6 from utils import load_data, plot_numeric_distributions
7 import pandas as pd
8
9 # %% [markdown]
10 # # Loading the data
11
12 # %%

```

```

13     file_name = "Small Data future AF"
14     data_path = f"../data/raw/{file_name}.csv"
15
16     # Read the CSV file
17     df = load_data(data_path)
18
19     # %% [markdown]
20     # ## Cleaning and renaming
21
22     # %%
23     df = df.rename(
24         columns={
25             "stock__share_type__company__industry__industry_group__"
26             "sector__name": (
27                 "sector"
28             ),
29             "stock__share_type__company__country__region": "region",
30             "market_cap_usd": "market_cap",
31             "stock__share_type__company__name": "company",
32             "stock__ticker": "ticker",
33         }
34     )
35
36     df["dividend"] = df["dividend"].fillna(0.0)
37     df["date"] = pd.to_datetime(df["date"])
38     df_no_nan = df.dropna()
39
40     # %% [markdown]
41     # # Analyzing
42
43     # %%
44     export_path = f"../data/cleaned/basic/plots/{file_name}.png"
45     plot_numeric_distributions(df, file_path=export_path)
46
47     # %% [markdown]
48     # # Export cleaned data
49
50     # %%
51     # Export the cleaned dataset to CSV
52     cleaned_data_path = f"../data/cleaned/basic/{file_name}.csv"
53     df.to_csv(cleaned_data_path, index=False)
54
55     print(f"Cleaned data exported to: {cleaned_data_path}")

```

**Listing 7.9:** Basic Data Cleaning

## 7.5.2 Outlier Detection

**Listing 7.10:** Outlier Detection

## 7.6 Extra Images

## 7.6.1 Data

### Columns and variables

Response Variables	Explanatory Variables				Time	Extra
Profit 1 month	Scores	Dummies		Size	Date	Company Name
Profit 3 months	Quality	Region	Sector	Market cap		Stock Ticker
Profit 6 months	Growth	USA and Canada	Communication Services	Volume		
Profit 1 year	Value	Europe	Consumer Discretionary			
Profit 2 years	Dividend	Latin America	Consumer Staples			
Profit 5 years		Asia and Oceania	Energy			
		Africa	Financials			
			Health Care			
			Industrials			
			Information Technology			
			Materials			
			Real State			
			Utilities			
6	4	1	1	2	1	2

Figure 7.1: Number of columns per dataset

## 7.6.2 Distributions

Here we have all of the not shown images for further analysis, if needed.

### Scores




Figure 7.2: Big Data Past - Scores










Figure 7.3: Big Data Past - Scores Deep

**Figure 7.4:** Small Data Future - Scores**Figure 7.5:** Small Data Future - Scores Deep

## Dummies

**Figure 7.6:** Big Data Future - Dummies**Figure 7.7:** Big Data Past - Dummies

## Profits

**Figure 7.8:** Big Data Future - Profits Boxplot**Figure 7.9:** Big Data Future - Profits

### 7.6.3 Correlations

#### USA and Canada




Figure 7.10: Small Data Future USA




Figure 7.11: Small Data Future USA - Multi

#### Europe




Figure 7.12: Small Data Future EU




Figure 7.13: Small Data Future EU Multi

#### Latin America




Figure 7.14: Small Data Future LAT




Figure 7.15: Small Data Future LAT - Multi

## Africa




Figure 7.16: Small Data Future AF




Figure 7.17: Small Data Future AF - Multi

## Asia and Oceania




Figure 7.18: Small Data Future AS




Figure 7.19: Small Data Future AS - Multi